

**CS186: Introduction to Database Systems**

Joe Hellerstein

Fall 2007

### Queries for Today

- What?
- Why?
- Who?
- How?
- For instance?

### What: Database Systems Then

### What: Database Systems Today

### What: Database Systems Today

### What: Database Systems Today

Account	Account Number	Available Balance
CHECKING	111-2000xx	\$7,299.46
SAVINGS	557-2911xx	\$188.46
SAVINGS	111-1535xx	\$1,252.65
<b>Total</b>		<b>\$8,739.57</b>

Account	Account Number	Total Account Value
BROKERAGE	W774xxxx	\$15,866.56
<b>Total</b>		<b>\$15,866.56</b>

Account	Account Number	Outstanding Balance	Available Credit
MASTERCARD	5490-9600-0008-xxxx	\$1,631.79	\$6,668.21
<b>Total</b>		<b>\$1,631.79</b>	<b>\$6,668.21</b>

Account	Account Number	Outstanding Principle Balance
STUDENTLOAN	70004xxxx	\$5,000.00
<b>Total</b>		<b>\$5,000.00</b>

## What: Database Systems Today

The screenshot shows a Yahoo! Local search results page for 'Berkeley'. The main area displays a map of the city with numerous orange location markers. A sidebar on the left lists categories such as 'GET MAP AND DIRECTIONS', 'FIND A BUSINESS ON THE MAP', and 'Browse by Category' (e.g., Restaurants, Hotels, Services). The top navigation bar includes links for News, Mail, and Search.

## What: Database Systems Today

The screenshot shows the NCBI Map Viewer for the Homo sapiens genome. The main display shows the human genome's chromosomes (1-22 and X) with a vertical color bar indicating GC content. The interface includes a search bar, tabs for different genome components, and a 'BLAST search the human genome' link. The bottom of the screen shows a timeline of genome releases from September 2006 to present.

## So... What is a Database?

- We will be broad in our interpretation
- A Database:
  - A very large, integrated collection of data.
- Typically models a real-world "enterprise"
  - Entities** (e.g., teams, games)
  - Relationships (e.g. **The A's are playing in the World Series**)
- Might surprise you how flexible this is
  - Web search:
    - Entities: words, documents
    - Relationships: **word in document**, **document links to document**.
  - P2P filesharing:
    - Entities: words, filenames, hosts
    - Relationships: **word in filename**, **file available at host**

## What is a Database Management System?

- A **Database Management System (DBMS)** is:
  - A software system designed to store, manage, and facilitate access to databases.
- Typically this term used narrowly
  - Relational databases with transactions
    - E.g. Oracle, DB2, SQL Server
  - Mostly because they predate other large repositories
    - Also because of technical richness
  - When we say **DBMS** in this class we will usually follow this convention
    - But keep an open mind about applying the ideas!

## What: Is the WWW a DBMS?

- That's a complicated question!
- The "surface web": docs and search
  - Crawler **indexes** pages on the web
  - Keyword-based **search** for pages
- Web-cache SW at Google/Yahoo is a kind of DBMS
- Notes
  - source data is mostly "prose": **unstructured** and **untyped**
  - public interface is **search only**:
    - can't modify the data
    - can't get summaries, complex combinations of data
  - few guarantees** provided for freshness of data, consistency across data items, fault tolerance, ...

## What: "Search" vs. Query

- Try **actors who donated to presidential candidates** in your favorite search engine.
- Now try **engineers who donated to presidential candidates**

The screenshot shows a Google search results page for the query 'engineers who donated to presidential candidates'. The search bar at the top contains the query. Below the search bar, there are several search results, each with a snippet of text and a link to a news article or website. The results are diverse, ranging from political news to engineering-related news stories.

 What: A "Database Query" Approach



Presented by the Federal Election Commission

Individuals Who Gave To: **KERRY, JOHN F.**

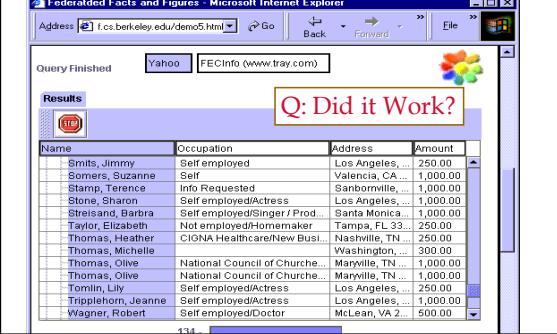
Sorted By Transaction Type Then Last Name

Committee(s) Used In This Query:

- KERRY-EDWARDS 2004 INC.**
- KERRY-EDWARDS 2004 INC. GENERAL ELECTION LEGAL ACCOUNTING COMPLIANCE FUND**
- JOHN KERRY FOR PRESIDENT, INC.**
- VETERANS FOR JUSTICE**

The query you have chosen matched 222599 individual contributions.

 "Yahoo Actors" JOIN "FECInfo"  
(Courtesy of the Telegraph research group @Berkeley)



Name	Occupation	Address	Amount
Smits, Jimmy	Self employed	Los Angeles, CA ...	250.00
Somers, Suzanne	Self	Valencia, CA ...	1,000.00
Stamp, Terence	Info Requested	Sanbonville, CA ...	1,000.00
Stone, Sharon	Self employed/Actress	Los Angeles, CA ...	1,000.00
Streisand, Barbra	Self employed/Singer / Prod...	Santa Monica, CA ...	1,000.00
Thorne, Heather	Not employed/Homemaker	Madison, WI ...	250.00
Thomas, Michelle	CIONA Healthcare New Busi...	Nashville, TN ...	250.00
Thomas, Olive	National Council of Churche...	Marshall, TN ...	300.00
Thomas, Olive	National Council of Churche...	Marshall, TN ...	1,000.00
Tomlin, Lily	Self employed/Actress	Los Angeles, CA ...	250.00
Tripplehorn, Jeanne	Self employed/Actress	Los Angeles, CA ...	1,000.00
Wagner, Robert	Self employed/Doctor	McLean, VA ...	500.00

 What: Is a File System a DBMS?

- Thought Experiment 1:
  - You and your project partner are editing the same file.
  - You both save it at the same time.
  - Whose changes survive?

**A) Yours   B) Partner's   C) Both   D) Neither   E) ???**

- Thought Experiment 2:
  - You're updating a file.
  - The power goes out.
  - Which changes survive?

**A) All   B) None   C) All Since Last Save   D) ???**

 What: Is a File System a DBMS?

- Thought Experiment 1:
  - You and your project partner are editing the same file.
  - You both save it at the same time.
  - Whose changes survive?

**A) Your changes   B) Partner's changes   C) Both changes   D) Neither changes   E) ???**

**A: Very, very carefully!!**

**A) All   B) None   C) All Since Last Save   D) ???**

 OS Support for Data Management

- Data can be stored in RAM
  - this is what every programming language offers!
  - RAM is fast, and random access
  - Isn't this heaven?
- Every OS includes a File System
  - manages *files* on a magnetic disk
  - allows *open, read, seek, close* on a file
  - allows protections to be set on a file
  - drawbacks relative to RAM?

 Database Management Systems

- What more could we want than a file system?
  - Simple, efficient *ad hoc*<sup>1</sup> queries
  - concurrency control
  - recovery
  - benefits of good data modeling
- S.M.O.P.<sup>2</sup>? Not really...
  - as we'll see this semester
  - in fact, the OS often gets in the way!

<sup>1</sup>**ad hoc:** formed or used for specific or immediate problems or needs  
<sup>2</sup>**SMOP:** Small Matter Of Programming



## Current Commercial Outlook

- Relational DBs a major part of the software industry
  - Oracle, IBM, Microsoft, HP, Teradata, Sybase, ...
- Open Source coming on strong
  - Relational: MySQL, PostgreSQL, Apache Derby, SQLite, Ingres, ...
  - text-Search: Lucene, Ferret, ...
- Well-known benchmarks (TPC, TREC)
- Tons of applications, related industries
  - Alphabet soup!
- Related database technologies have niches
  - P2P, XML repositories, etc.



## What systems will we cover?

- We will try to be broad and touch upon
  - Relational **DBMS** (e.g. Oracle, SQL Server, DB2, Postgres)
  - Document **search engines** (e.g. Google, Yahoo! Search, Lucene, Ferret)
- Ground things in relevant applications



## Quiz Question

- Name some widely-used applications



## Why take this class?

- Database systems are at the core of CS
- They are incredibly important to society
- The topic is intellectually rich
- A capstone course for undergrad
- ~~E. It isn't that much work~~
- Looks good on your resume

Let's spend a little time on each of these



## Why take this class?

### A. Database systems are the core of CS

- Shift from computation to information
  - True in corporate computing for years
  - Web made this clear for "the rest of us" by the end of 90's
  - Increasingly true of scientific computing
- Need for DB technology has exploded in the last years
  - **Corporate:** retail swipe/clickstreams, "customer relationship mgmt", "supply chain mgmt", "data warehouses", etc.
  - **Web:** not just "documents". Search engines, maps, e-commerce, blogs, wikis, social networks. Web 2.0.
  - **Scientific:** digital libraries, genomics, satellite imagery, physical sensors, simulation data
  - **Personal:** Music, photo, & video libraries. Email archives. File contents ("desktop search").



## Why take this class?

### B. DBs are incredibly important to society

- "Knowledge is power." -- Sir Francis Bacon
- "With great power comes great responsibility." -- Spiderman's Uncle Ben



- Policy-makers should understand technological possibilities.
- Informed Technologists needed in public discourse on usage.



## Why take this class?

### C. The topic is intellectually rich.

- representing information
  - data modeling
- languages and systems for querying data
  - complex queries & query semantics\*
  - over massive data sets
- concurrency control for data manipulation
  - controlling concurrent access
  - ensuring transactional semantics
- reliable data storage
  - maintain data semantics even if you pull the plug

\* semantics: the meaning or relationship of meanings of a sign or set of signs



## Why take this class?

### D. The course is a capstone.

- We will see
  - Algorithms and cost analyses
  - System architecture and implementation
  - Resource management and scheduling
  - Language design, semantics and optimization
  - AI topics including logic and planning
  - Statistical modeling of data



## Why take this class?

### E. It isn't that much work.

- Bad news: It is a fair bit of work.
  - varies from year to year
- Good news: the course is front loaded
  - Most of the hard work is in the first half of the semester
  - Load balanced with most other classes



## Why take this class?

### F. Looks good on my resume.

- Yes, but why? This is not a course for:
  - Oracle administrators
  - IBM DB2 engine developers
  - Though it's useful for both!
- It is a course for well-educated computer scientists
  - Database system concepts and techniques increasingly used "outside the box"
    - Ask your friends at Microsoft, Yahoo!, Google, Apple, etc.
    - Actually, they may or may not realize it!
  - A rich understanding of these issues is a basic and (un?)fortunately unusual skill.



## Who?

- Instructor
  - Prof. Joe Hellerstein
  - cs186profs@db.cs.berkeley.edu
- TAs
  - David Chu
  - Eirinaios Michelakis



## How? Workload

- Projects with a "real world" focus:
  - Modify the internals of a "real" open-source database system: PostgreSQL
    - Serious C system hacking in a ~500KLoc codebase
    - Measure the benefits of our changes
  - Build web-based applications
    - Using Ruby on Rails, PostgreSQL, Ferret text search
- Other homework assignments and/or quizzes
- Exams – 1 Midterm & 1 Final
- Projects to be done in groups of 2
  - Pick your partner ASAP
- The course is "front-loaded"
  - most of the hard work is in the first half



## How? Administrivia

- <http://inst.eecs.berkeley.edu/~cs186>
- Office Hours:
  - JMH:
    - Tues 12:30-1:30
    - Thurs 11:00-12:00 (in 685 Soda)
  - DC: Mon 1-2
  - EM: Fri 11-12 (locations TBA)
- Discussion Sections **WILL** meet this week



## How? Administrivia, cont.

- textbook
  - *Database Management Systems, 3rd Edition*
    - Ramakrishnan and Gehrke
  - *Agile Web Development with Rails, 2nd edition*
    - e-book is fine (better?)
  - *Programming Ruby, 2nd edition*
    - Free online
- Grading, hand-in policies, etc. will be on Web Page
- Cheating policy: zero tolerance
  - We have the technology...



## How? Administrivia, cont.

- Team Projects
  - Teams of 2
  - Think about this now!
- Class bulletin board - ucb.class.cs186
  - read it regularly and post questions/comments.
  - mail broadcast to all TAs will not be answered
  - mail to the cs186 course account will not be answered
- Class Blog for announcements



## Agenda for the rest of today

- A “free tasting” of central concepts in DB field:
  - queries and search
  - data independence
  - transactions
- Next Time
  - the Relational data model
  - object-relational mapping using Ruby on Rails
- Today’s lecture is from Chapter 1 in R&G
- Read Chapter 2 for next class.



## Describing Data: Data Models

- A *data model* is a collection of concepts for describing data.
- A *schema* is a description of a particular collection of data, using a given data model.
- The *relational model of data* is the most widely used model today.
  - Main concept: *relation*, basically a table with rows and columns.
  - Every relation has a *schema*, which describes the columns, or fields.



## Example: University Database

- Schema:
  - Students(sid text,  
          name text,  
          login text,  
          age integer,  
          gpa float)
  - Courses(cid text,  
          cname text,  
          credits integer)
  - Enrolled(sid text,  
          cid text,  
          grade text)

### Levels of Abstraction

**Users**

- Views describe how users see the data.
- Conceptual schema defines logical structure
- Physical schema describes the files and indexes used.

```

graph TD
    View1[View 1] --> CS[Conceptual Schema]
    View2[View 2] --> CS
    View3[View 3] --> CS
    CS --> PS[Physical Schema]
    PS --> DB[DB]
    
```

The diagram illustrates the levels of abstraction. At the top, three icons represent different users: a duck, a person at a computer, and a cat. Below them are three boxes labeled "View 1", "View 2", and "View 3". Arrows point from each view box to a central box labeled "Conceptual Schema". An arrow points from the "Conceptual Schema" box down to a box labeled "Physical Schema". Finally, an arrow points from the "Physical Schema" box down to a cylinder labeled "DB".

### Example: University Database

- Conceptual schema:
  - Students(sid text, name text, login text, age integer, gpa float)
  - Courses(cid text, cname text, credits integer)
  - Enrolled(sid text, cid text, grade text)
- Physical schema:
  - Relations stored as unordered files.
  - Index on first column of Students.
- External Schema (View):
  - Course\_info(cid text, enrollment integer)

### Data Independence

- Applications insulated from how data is structured and stored.
- Logical data independence: Protection from changes in *logical* structure of data.
- Physical data independence: Protection from changes in *physical* structure of data.
- Q: Why is this particularly important for DBMS?

Because databases and their associated applications persist.

### Hellerstein's Inequality

$$\frac{dapp}{dt} \ll \frac{denv}{dt}$$

### Agenda ...

- A “free tasting” of central concepts in DB field:
  - queries (vs. search)
  - data independence
  - – transactions

### Concurrent execution of user programs

- Why?
  - Utilize CPU while waiting for disk I/O
    - (database programs make heavy use of disk)
  - Avoid short programs waiting behind long ones
    - e.g. ATM withdrawal while bank manager sums balance across all accounts



## Concurrent execution

- Interleaving actions of different programs: trouble!

**Example:**

- Bill transfers \$100 from savings to checking  
 $Savings -= 100; Checking += 100$
- Meanwhile, Bill's wife requests account info.

**Bad interleaving:**

- Savings -= 100
- Print balances
- Checking += 100

– Printout is missing \$100 !



## Concurrency Control

- DBMS ensures such problems don't arise
- Users can pretend they are using a single-user system. (called "**Isolation**")  
– Thank goodness!



## Key concept: Transaction

- an **atomic sequence** of database actions (reads/writes)
- takes DB from one **consistent state** to another



## Example



- Here, **consistency** is based on our knowledge of banking "semantics"
- In general, up to writer of transaction to ensure transaction preserves consistency
- DBMS provides (limited) automatic enforcement, via **integrity constraints**  
– e.g., balances must be  $\geq 0$



## Concurrent transactions

- Goal: **execute xacts {T1, T2, ... Tn}, and ensure a consistent outcome**
- **One option:** "serial" schedule (**one after another**)
- **Better:** allow interleaving of xact actions, as long as outcome is equivalent to **some serial schedule**



## Possible Enforcement Methods

- Optimistic: **permit arbitrary interleaving, then check equivalence to serial sched.**
- Pessimistic: **xacts set locks on data objects, such that illegal interleaving is impossible**



## Locking example

- T1 (Bill): *Savings -= 100; Checking += 100*
- T2 (Bill's wife): *Print(Checking); Print(Savings)*
- T1 and T2 both lock Savings and Checking objects
- If T1 locks Savings & Checking first, T2 must wait



## A wrinkle ...

- T1 (Bill): *Savings -= 100; Checking += 100*
- T2 (Bill's wife): *Print(Checking); Print(Savings)*

Suppose:

1. T1 locks Savings
  2. T2 locks Checking
- Now neither transaction can proceed!
  - called "deadlock"
  - DBMS will abort and restart one of T1 and T2
  - Need "undo" mechanism that preserves consistency
  - Undo mechanism also necessary if system crashes between "Savings -= 100" and "Checking += 100" ...



## Ensuring Transaction Properties

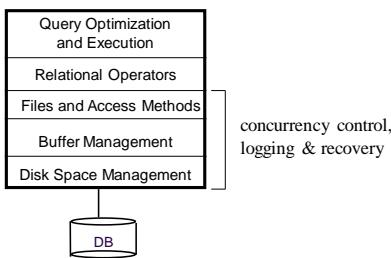
- DBMS ensures:
  - **atomicity** even if xact aborted (due to deadlock, system crash, ...)
  - **durability** of committed xacts, even if system crashes.
- Idea: Keep a *log* of all actions carried out by the DBMS:
  - Record all DB modifications in log, *before* they are executed
  - To abort a xact, undo logged actions in reverse order
  - If system crashes, must:
    - 1) undo partially executed xacts (ensures **atomicity**)
    - 2) redo committed xacts (ensures **durability**)
  - trickier than it sounds!



## Architecture of a DBMS ...



## Typical DBMS architecture



## Advantages of a Traditional DBMS

- Data independence
- Efficient data access
- Data integrity & security
- Data administration
- Concurrent access, crash recovery
- Reduced application development time
- So why not use them always?
  - Expensive/complicated to set up & maintain
  - This cost & complexity must be offset by need
  - General-purpose, not suited for special-purpose tasks (e.g. text search!)



## Databases make these folks happy ...

- Web & enterprise app developers
- Computing infrastructure providers
- DBMS vendors, programmers
  - Oracle, IBM, MS ...
- End users in *many* fields
  - Business, education, science, ...
- Database administrators (DBAs)



...must understand how a DBMS works



## Summary

- Relational DBMS: maintain/query structured data
  - broadly applicable
  - can manipulate data and exploit *semantics*
  - recovery from system crashes
  - concurrent access
  - robust application development and *evolution*
  - data integrity and security
- Text search engine
  - similar to relations underneath
  - many “application-specific” smarts



## Summary, cont

- Levels of abstraction & data independence.
  - Hellerstein’s inequality
  - classic idea, resonates in the most modern SW
- Goals of the course
  - 1) How to be a sophisticated user of database technology
  - 2) What goes on inside a DBMS and search engine
  - 3) How to architect data-intensive systems