

External sorting

R & G – Chapter 13

Guest Lecturer:

Brian Cooper
Yahoo! Research



A little bit about Y!

- Yahoo! is the most visited website in the world
 - Sorry Google
 - 500 million unique visitors per month
 - 74 percent of U.S. users use Y! (per month)
 - 13 percent of U.S. users' online time is on Y!



Why sort?

Sort by: Relevance Distance			Showing 1 to 15 of 200
Business Name	Address	Rating	Reviews
King Pin Doughnuts (510) 843-6089 See reviews on Local	2521 Durant Ave #A Berkeley, CA Map	0.2	
Roath's Bagels (510) 849-9551 See reviews on Local	2341 Telegraph Ave Berkeley, CA Map	0.2	
Dream Fruit Donuts (510) 849-8471 See reviews on Local	2637 Ashby Ave Berkeley, CA Map	1.0	
Roath's Bagels (510) 854-0944 See reviews on Local	3170 College Ave Berkeley, CA Map	1.4	
All Star Donut (510) 868-6078 See reviews on Local	1255 University Ave Berkeley, CA Map	1.5	
Roath's Bagels (510) 525-4447 See reviews on Local	1883 Solano Ave Berkeley, CA Map	1.7	
Boogie Vibe Bagel Boy (510) 524-3194 See reviews on Local	1281 Gilman St Albany, CA Map	1.8	
Boogie Vibe Bagel Boy (510) 527-8272 See reviews on Local	1218 Santa Fe Ave Albany, CA Map	1.8	
Berkeley Donut Shop (510) 853-9044 See reviews on Local	3043 San Pablo Ave Berkeley, CA Map	2.0	
Happy Donuts (510) 524-8816 See reviews on Local	1041 Gilman St Berkeley, CA Map	2.1	

amazon.com


"Toy" > Toys & Games

Showing 1 - 24 of 260,516 Results | Page: 1 2 3 ... | Next > Sort by: Price: High to Low

<p>Steffi Germany: Giant Studio Elephant: Overall Size ~ 210cm high (82.68") Buy now: \$22,000.00 Buy now: \$16,000.00 Usually ships in 2 to 3 weeks</p>	<p>Miss Megan Modular Playground 3.5 Inch Poles Buy now: \$12,922.00 Buy now: \$12,922.00 Usually ships in 2 to 3 weeks</p>	<p>Heade LX200 GPS 16 in. UHTC SCT with Super Field Tripod Buy now: \$10,988.71 Buy now: \$10,988.71 In Stock</p>
<p>Apollo 17 Astronaut Space Suit Replica Currently unavailable</p>	<p>Meade 14" f/8 RCT Advanced Ritchey-Chretien Telescope with UHTC Tripod - 1408-45-01 Buy now: \$44,494.00 Buy now: \$9,599.99 2 Used & New from \$9,593.71 In Stock</p>	<p>Lizard Thumb Piece Entry Way Lock Set - ETS2418 - Thumbgrip Handsets Currently unavailable</p>

Why sort?

- Users usually want data sorted
- Sorting is first step in bulk-loading a B+ tree
- Sorting useful for eliminating duplicates
- Sort-merge join algorithm involves sorting



So?

- Don't we know how to sort?
 - Quicksort
 - Mergesort
 - Heapsort
 - Selection sort
 - Insertion sort
 - Radix sort
 - Bubble sort
 - Etc.
- Why don't these work for databases?

Key problem in database sorting




4 GB: \$300




480 GB: \$300

- How to sort data that does not fit in memory?

Example: merge sort



Example: merge sort




Isn't that good enough?


- Consider a file with N records
- Merge sort is $O(N \lg N)$ comparisons
- We want to minimize disk I/Os
 - Don't want to go to disk $O(N \lg N)$ times!
- Key insight: **sort based on pages, not records**
 - Read whole pages into RAM, not individual records
 - Do some in-memory processing
 - Write processed blocks out to disk
 - Repeat

2-way sort


- Pass 0: sort each page



- Pass 1: merge two pages into one run



- Pass 2: merge two runs into one run



- ...
- Sorted!

What did that cost us?

- P pages in the file
- Each pass: read and wrote P pages
- How many passes?
 - Pass 0
 - Pass 1: went from P pages to $P/2$ runs
 - Pass 2: went from $P/2$ runs to $P/4$ runs
 - ...
 - Total number of passes: $\lceil \log_2 P \rceil + 1$
- Total cost: $2P * (\lceil \log_2 P \rceil + 1)$

What did that cost us?

- Why is this better than plain old merge sort?
 - $N \gg P$
 - So $O(N \lg N) \gg O(P \lg P)$
- Example:
 - 1,000,000 record file
 - 8 KB pages
 - 100 byte records
 - = 80 records per page
 - = 12,500 pages
 - Plain merge sort: 41,863,137 disk I/O's
 - 2-way external merge sort: 365,241 disk I/O's
 - 4.8 days versus 1 hour

Can we do better?


- 2-way merge sort only uses 3 memory buffers
 - Two buffers to hold input records
 - One buffer to hold output records
 - When that buffer fills up, flush to disk
- Usually we have a lot more memory than that
 - Set aside 100 MB for sort scratch space = 12,800 buffer pages
- Idea: read as much data into memory as possible each pass
 - Thus reducing the number of passes
 - Recall total cost:

$2P * \text{Passes}$


External merge sort

- Assign B input buffers and 1 output buffer
- Pass 0: Read in runs of B pages, sort, write to disk
- Pass 1: Merge B runs into one
 - For each run, read one block
 - When a block is used up, read next block of run
- Pass 2: Merge B runs into one
- ...
- Sorted!


Example




Example




Example




Example




Example




Example




Example




Example




Example




Example




Example




Example




Example




Example




Example



Example



Example



What did that cost us?

- P pages in file, B buffer pages in RAM
- P/B runs of size B
- Each pass: read and write P pages
- How many passes?
 - $\lceil \log_{B-1} \lceil P/B \rceil \rceil + 1$
- Total cost: $2P * \lceil \log_{B-1} \lceil P/B \rceil \rceil + 1$

Example

- 1,000,000 records in 12,500 pages
- Use 10 buffer pages in memory
- 4 passes
- 100,000 disk I/Os
 - 17 minutes versus 1 hour for 2-way sort

Can I do two passes?

- Pass 0: sort runs
- Pass 1: merge runs
- Given B buffers
- Need:
 - No more than B-1 runs
 - Each run no longer than B pages
- Can do two passes if $P \leq B * (B-1)$
- Question: what's the largest file we can sort in three passes? N passes?

Make I/Os faster


- Cost = I/Os is a simplification
 - Sequential I/Os are cheaper than random I/Os
- Read blocks of pages at a time
 - X = Blocking factor
 - B = buffer pages
 - $(B/X - X)$ input "buffer blocks", one output "buffer block"
- Result
 - Fewer runs merged per pass = more passes
 - Less time per I/O = quicker passes
 - Tradeoff!
 - Maximize total sort time by choosing X given B, P and I/O latencies

Overlap computation and I/O

- Problem: CPU must wait for I/O
 - Suppose I need to read a new block
 - Stop merging
 - Initiate I/O
 - Wait
 - Complete I/O
 - Resume merging


Solution: double buffering

- Keep a second set of buffers
 - Process one set while waiting for disk I/O to fill the other set




Solution: double buffering

- Keep a second set of buffers
 - Process one set while waiting for disk I/O to fill the other set




Solution: double buffering

- Keep a second set of buffers
 - Process one set while waiting for disk I/O to fill the other set




Solution: double buffering

- Keep a second set of buffers
 - Process one set while waiting for disk I/O to fill the other set




Solution: double buffering

- Keep a second set of buffers
 - Process one set while waiting for disk I/O to fill the other set



Solution: double buffering

- Keep a second set of buffers
 - Process one set while waiting for disk I/O to fill the other set




What if the data is already sorted?

- Yay!
- Often this happens because of a B+ tree index
 - Leaf level of a B+ tree has all records in sorted order
 - Two possibilities: B+ tree is **clustered** or **unclustered**


Clustered B+ tree

Sweep through leaf layer, reading data blocks in order



Clustered B+ tree

Sweep through leaf layer, reading leaf blocks in order



What did that cost us?


- Traverse B+ tree to left-most leaf page
- Read all leaf pages
 - For each leaf page, read data pages
- Data not in B+ tree:
 - Height + Width + Data pages
- Data in B+ tree:
 - Height + Width

Example

- 1,000,000 records, 12,500 data pages
- Assume keys are 10 bytes, disk pointers are 8 bytes
 - So ≈ 300 entries per 8 KB B+ tree page (if two-thirds full)
- Data not in B+ tree
 - 12,500 entries needed = 42 leaf pages
 - Two level B+ tree
 - Total cost: $1 + 42 + 12,500 = 12,543$ I/Os
 - 2 minutes versus 17 minutes for external merge sort
- Data in B+ tree
 - Three level B+ tree, 12,500 leaf pages
 - Total cost: $2 + 12,500 = 12,502$ I/Os
 - Also about 2 minutes

What if the B+ tree is unclustered?

- We know the proper sort order of the data
- But retrieving the data is hard!



What if the B+ tree is unclustered?

- Result is that in the worst case, may need one disk I/O per record
 - Even though we know the sort order!
- Usually external merge sort is better in these cases
 - Unless all you need is the set of keys

Summary

- Sorting is very important
- Basic algorithms not sufficient
 - Assume memory access free, CPU is costly
 - In databases, memory (e.g. disk) access is costly, CPU is (almost free)
- Try to minimize disk accesses
 - 2-way sort: read and write records in blocks
 - External merge sort: fill up as much memory as possible
 - Blocked I/O: try to do sequential I/O
 - Double buffering: read and compute at the same time
 - Clustering B+ tree: the data is already sorted. Hooray!
 - Unclustered B+ tree: no help at all

