


Tree-Structured Indexes

R & G Chapter 10

"If I had eight hours to chop down a tree, I'd spend six sharpening my ax."

Abraham Lincoln



Review: Files, Pages, Records

- Abstraction of stored data is "files" with "pages" of "records".
 - Records live on *pages*
 - Physical Record ID (RID) = <page#, slot#>
- Variable length* data requires more sophisticated structures for records and pages. (why?)
 - Fields in Records: offset array in header
 - Records on Pages: Slotted pages w/internal offsets & free space area
- Often best to be "lazy" about issues such as free space management, exact ordering, etc. (why?)
- Files can be unordered (heap), sorted, or kinda sorted (i.e., "clustered") on a *search key*.
 - Tradeoffs are update/maintenance cost vs. speed of accesses via the search key.
 - Files can be clustered (sorted) at most one way.
- Indexes can be used to speed up many kinds of accesses. (i.e., "access paths")



Tree-Structured Indexes: Introduction

- Selections of form field <op> constant
- Equality** selections (op is =)
 - Either "tree" or "hash" indexes help here.
- Range** selections (op is one of <, >, <=, >=, BETWEEN)
 - "Hash" indexes *don't* work for these.
- More complex selections** (e.g. spatial containment)
 - There are fancier trees that can do this... more on this soon!
- Tree-structured indexing techniques support both *range selections* and *equality selections*.
- ISAM**: static structure; early index technology.
- B+ tree**: dynamic, adjusts gracefully under inserts and deletes.
- ISAM = Indexed Sequential Access Method



A Note of Caution

- ISAM is an old-fashioned idea
 - B+-trees are usually better, as we'll see
 - Though not *always*
- But, it's a good place to start
 - Simpler than B+-tree, but many of the same ideas
- Upshot
 - Don't** brag about being an ISAM expert on your resume
 - Do** understand how they work, and tradeoffs with B+-trees



Range Searches

- "Find all students with gpa > 3.0"
 - If data is in sorted file, do binary search to find first such student, then scan to find others.
 - Cost of binary search in a database can be quite high. Q: Why???
- Simple idea: Create an 'index' file.




☒ Can do binary search on (smaller) index file!



ISAM

- Index file may still be quite large. But we can apply the idea repeatedly!




☒ Leaf pages contain *data entries*.



Example ISAM Tree

- **Index entries:** <search key value, page id> they direct search for data entries *in leaves*.
- Example where each node can hold 2 entries;



ISAM is a STATIC Structure


- **File creation:** Leaf (data) pages allocated sequentially, sorted by search key; then index pages allocated, then overflow pgs.
- **Search:** Start at root; use key comparisons to go to leaf. Cost = $\log_F N$; $F = \# \text{ entries/pg (i.e., fanout)}$, $N = \# \text{ leaf pgs}$
 - no need for 'next-leaf-page' pointers. (Why?)
- **Insert:** Find leaf that data entry belongs to, and put it there. Overflow page if necessary.
- **Delete:** Find and remove from leaf; if empty page, de-allocate.

Data Pages
Index Pages
Overflow pages


Static tree structure: inserts/deletes affect only leaf pages.



Example: Insert 23*, 48*, 41*, 42*



... then Deleting 42*, 51*, 97*



☒ Note that 51* appears in index levels, but not in leaf!




ISAM ---- Issues?

- Pros
 - ????
- Cons
 - ????



B+ Tree: The Most Widely Used Index

- Insert/delete at $\log_F N$ cost; keep tree **height-balanced**.
 - $F = \text{fanout}$, $N = \# \text{ leaf pages}$
- Minimum 50% occupancy (except for root). Each node contains m entries where $d \leq m \leq 2d$ entries. "d" is called the order of the tree.
- Supports equality and range-searches efficiently.
- As in ISAM, all searches go from root to leaves, but structure is **dynamic**.






Example B+ Tree

- Search begins at root, and key comparisons direct it to a leaf (as in ISAM).
- Search for 5*, 15*, all data entries $\geq 24^*$

...



✉ Based on the search for 15*, we know it is not in the tree!



B+ Trees in Practice

- Typical order: 100. Typical fill-factor: 67%.
 - average fanout = 133
- Typical capacities:
 - Height 2: $133^2 = 17,689$ entries
 - Height 3: $133^3 = 2,352,637$ entries
- Can often hold top levels in buffer pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 Mbytes




Inserting a Data Entry into a B+ Tree

- Find correct leaf L .
- Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must *split* L (into L and a new node $L2$)
 - Redistribute entries evenly, *copy up* middle key.
 - Insert index entry pointing to $L2$ *into parent* of L .
- This can happen recursively
 - To *split index node*, redistribute entries evenly, but *push up* middle key. (Contrast with leaf splits.)
- Splits "grow" tree; root split increases height.
 - Tree growth: gets *wider* or *one level taller at top*.



Example B+ Tree - Inserting 8*



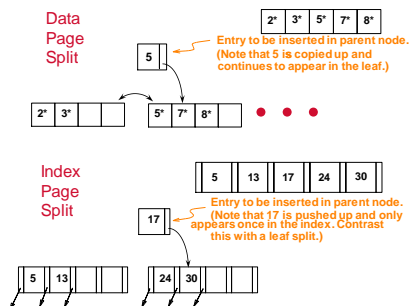
❖ Notice that root was split, leading to increase in height.

❖ In this example, we can avoid split by re-distributing entries; however, this is usually not done in practice.



Data vs. Index Page Split (from previous example of inserting "8*")

- Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- Note difference between *copy-up* and *push-up*; be sure you understand the reasons for this.



Deleting a Data Entry from a B+ Tree

- Start at root, find leaf L where entry belongs.
- Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only 0-1 entries,
 - Try to re-distribute, borrowing from *sibling* (adjacent node with same parent as L).
 - If re-distribution fails, *merge* L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- Merge could propagate to root, decreasing height.

In practice, many systems do not worry about ensuring half-full pages. Just let page slowly go empty; if it's truly empty, just delete from tree and leave unbalanced.



Deleting a Data Entry from a B+ Tree

- Start at root, find leaf L where entry belongs.
- Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only **d-1** entries,
 - Try to **re-distribute**, borrowing from *sibling* (adjacent node with same parent as L).
 - If re-distribution fails, **merge** L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- Merge could propagate to root, decreasing height.



Prefix Key Compression

- Important to increase fan-out. (Why?)
- Key values in index entries only 'direct traffic'; can often compress them.
 - E.g., If we have adjacent index entries with search key values *Dannon Yogurt*, *David Smith* and *Devarakonda Murthy*, we can abbreviate *David Smith* to *Dav*. (The other keys can be compressed too ...)
 - Is this correct? Not quite! What if there is a data entry *Davey Jones*? (Can only compress *David Smith* to *Dav*)
 - In general, while compressing, must leave each index entry greater than every key value (in any subtree) to its left.
- Insert/delete must be suitably modified.




Suffix Key Compression

- If many index entries share a common prefix
 - E.g. MacDonald, MacEnroe, MacFeeley
 - Store the common prefix "Mac" at a well known location on the page, use suffixes as split keys
- Particularly useful for composite keys
 - Why?



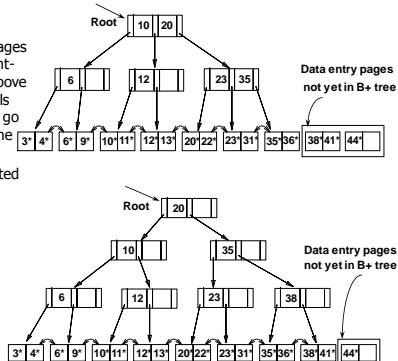
Bulk Loading of a B+ Tree

- If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.
 - Also leads to minimal leaf utilization --- why?
- Bulk Loading** can be done much more efficiently.
- Initialization.** Sort all data entries, insert pointer to first (leaf) page in a new (root) page.



Bulk Loading (Contd.)

- Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits. (Split may go up right-most path to the root.)
- Much faster than repeated inserts.



Summary of Bulk Loading

- Option 1: multiple inserts.
 - Slow.
 - Does not give sequential storage of leaves.
- Option 2: **Bulk Loading**
 - Fewer I/Os during build.
 - Leaves will be stored sequentially (and linked, of course).
 - Can control "fill factor" on pages.



A Note on 'Order'

- Order (**d**) concept replaced by physical space criterion in practice ('at least half-full').
 - Index pages can typically hold many more entries than leaf pages.
 - Variable sized records and search keys mean different nodes will contain different numbers of entries.
 - Even with fixed length fields, multiple records with the same search key value (*duplicates*) can lead to variable-sized data entries (if we use Alternative (3)).
- Many real systems are even sloppier than this --- only reclaim space when a page is *completely* empty.



Summary

- Tree-structured indexes are ideal for range-searches, also good for equality searches.
- ISAM is a static structure.
 - Only leaf pages modified; overflow pages needed.
 - Overflow chains can degrade performance unless size of data set and data distribution stay constant.
- B+ tree is a dynamic structure.
 - Inserts/deletes leave tree height-balanced; $\log_F N$ cost.
 - High fanout (**F**) means depth rarely more than 3 or 4.
 - Almost always better than maintaining a sorted file.



Summary (Contd.)

- Typically, **67%** occupancy on average.
- Usually preferable to ISAM, modulo *locking* considerations; adjusts to growth gracefully.
- If data entries are data records, splits can change rids!
- Key compression increases fanout, reduces height.
- Bulk loading can be much faster than repeated inserts for creating a B+ tree on a large data set.
- Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.