

## Generalized Search Trees and Spatial Indexing

(Ramakrishnan 28.1, 28.2, 28.6)



### Other Search Trees

- **Question: Can B+-trees handle more complicated searches?**
  - A typical example: "*gpa* > 3.7 *and* *age* < 18"
  - Same thing: "all restaurants in downtown Berkeley"
  - Even fancier: "all pictures resembling </tmp/sunset.gif>"
  - (Easy: "all pictures identical to </tmp/sunset.gif>")
- **B+-trees exploit data order to do range search**
  - 1-d range search is not always what you want




### Search Trees in General

- **Lots of trees invented for multidimensional data**
  - e.g. R-trees, R\*-trees, hB-trees, UB-trees, X-trees, etc. etc.
- **New tree indexes for other kinds of data**
  - Image/video search, timeseries matching, DNA sequence matching, etc.
- **Many are "unordered" B+-trees, fancy keys**
  - In some ways, B+-tree is just a "special case" of these trees.




### Search Trees from 30,000 feet




### A "Disorderly" B+-tree

- **How to search a disorderly B+-tree?**

- Equality search is identical to traditional!
- Range search = traversing multiple paths
  - follow all pointers where key range overlaps query range.



### Another Disorderly B+-Tree



- **Insert 41???**
  - keys need not cover all possible values
- **Search for 77??**
  - keys may "overlap"



## Generalized Search Tree (GiST)

- **A disorderly B+-tree, with *user-defined* keys**
  - tree doesn't interpret keys.
  - "user" implements keys as an OO class, with methods that guide search, insert, delete, split, etc.
- **Structure: balanced tree of  $(p, ptr)$  pairs**
  - $p$  is an index key or "predicate"
  - $p$  holds for all data records below  $ptr$
  - need  $n$  keys for  $n$  pointers (unlike B+-tree)



## User-provided Key Methods

- **Search:**
  - **Consistent** $(E, q)$ :  $E.p \wedge q$ ? (no/maybe)
- **Generating new keys after splits:**
  - **Union** $(P)$ : new key that holds for all tuples in  $P$
- **Data organization:**
  - **Penalty** $(E_1, E_2)$ :  
"badness" of inserting  $E_2$  in subtree  $E_1$
  - **PickSplit** $(P)$ : split  $P$  into two groups of entries



## Search

- **General technique:**
  - Depth-First Search where **Consistent** is TRUE
- **Incremental algorithm:**
  - Maintain a search stack of  $\langle \text{page\_id}, \text{offset} \rangle$  pairs from root to current data entry
  - Each "get next" call moves the offset to right
    - When nothing left on page, pop stack, move right in parent (or recurse) and go down to leaf.



## Insert

- **descend tree along least increase in Penalty**
- **if there's room at leaf, insert there**
- **else split according to PickSplit**
  - do B-tree-style recursive splitting
- **propagate changes (recursively) using Union**
  - make sure all ancestor keys are consistent with inserted item

Q: what happened to B-tree "copy up" and "push up" logic?

Let's revisit with example of R-tree in a few slides




## Delete

- **find the entry via Search, and delete it**
- **propagate changes (recursively) using Union**
- **on underflow:**
  - reinsert stuff on page and delete page
  - why not borrow/merge a la B+-trees?



## RD-tree: GiST for Sets





## RD-trees

- **Logically, keys represent minimal supersets**
- **Queries: Contains, Intersects, Equals...**
- **Consistent( $E, q$ ):**
  - Varies slightly on query type ... e.g.  $E.p \cap q \neq \emptyset$
- **Union( $P$ ): set-union of keys**
- **Penalty( $E, F$ ):  $|E.p \cup F.p| - |E.p|$**
- **PickSplit( $P$ ): many possible algorithms**
  - One goal: minimize sum of cardinalities of 2 pages

Used in Postgres as an alternative to inverted indexes  
Note: need key compression here (like postings lists)




## R-tree: A GiST over 2-d data

- **(Invented ~10 years before GiST or RD-Tree)**
  - Also at Berkeley ;-)
- **Logically, keys represent "bounding boxes"**
- **Queries: Contains, Overlaps, Equals ... *bbox***
- **Consistent( $E, q$ ):**
  - Varies slightly on query type: e.g.  $E.p$  overlaps  $q$ ?
- **Union( $P$ ): bounding box of all entries**
- **Penalty( $E, F$ ):  $\text{size}(\text{Union}(\{E, F\})) - \text{size}(E)$**
- **PickSplit( $P$ ):**
  - goal: minimize sum of areas of the 2 pages



## R-trees, Slowly. Problem:

- **Given a collection of geometric objects (points, lines, polygons, ...)**
- **organize them on disk, to answer spatial queries (range, near neighbors, etc)**



R-tree slides from G. Kollios, BU



## R-trees

- **Main idea: a 2-D B-tree with overlapping keys!**
  - => guaranteed 50% utilization
  - => easier insertion/split algorithms.
  - (only deal with Minimum Bounding Rectangles - **MBRs**)




## R-tree Structure

- **A multi-way tree of disk blocks**
- **Index nodes and data (leaf) nodes**
- **All leaf nodes appear on the same level**
- **Every node contains between m and M entries**
- **The root node has at least 2 entries (children)**



## Example


- **eg., w/ fanout 4: group nearby rectangles to parent MBRs; each group -> disk page**






### Example

- $F=4$




### Example

- $F=4$




### R-trees - format of nodes

- $\{(MBR; \text{obj\_ptr})\}$  for leaf nodes




### R-trees - format of nodes


- $\{(MBR; \text{node\_ptr})\}$  for non-leaf nodes



### R-trees:Search



### R-trees:Search




## R-trees: Search

- **Main points:**

- every parent node completely covers its 'children'
- a child MBR may be covered by more than one parent - it is stored under ONLY ONE of them. (ie., no need for dup. elim.)
- a point query may follow multiple branches.
- everything works for **any(?)** dimensionality

## R-trees: Near Neighbor Search

- **Rather than a stack (for depth-first), maintain a priority queue of nodes to visit**
- **Upon visiting a page, load all the children to be traversed into the priority queue**
  - Priority = MinDistance




## R-tree: Incremental Near Neighbors

```
PQ.insert(root)
while (!PQ.empty)
  next = PQ.deletemin
  if (next is a data entry) return (next)
  else // next is a page in the index
    for each item i on next
      PQ.insert(i)
    end for
  end if
end while
```


## R-trees: Insertion

Insert X




## R-trees: Insertion

Insert Y



## R-trees: Insertion

- **Extend the parent MBR**




## R-trees: Insertion

- **How to find the next node to insert the new object?**
  - Penalty metric: At each level, find the entry that needs the least enlargement to include Y. Resolve ties using the area (smallest)
- **Other methods have been proposed**
  - E.g. perhaps useful to minimize the perimeter of MBRs too?


## R-trees: Insertion

- If node is full then **Split** : ex. Insert w



## R-trees: PickSplit


- **Split node P1: partition the MBRs into two groups.**



- Naïve: exhaustive
  - cost?
- quadratic split
- linear split


## R-tree PickSplit: idea

- pick two rectangles as 'seeds';
- assign each rectangle 'R' to the 'closest' 'seed'



## R-trees: Split

- pick two rectangles as 'seeds';
- assign each rectangle 'R' to the 'closest' 'seed';
- 'closest': the smallest increase in area



## R-trees: Split

- **How to pick Seeds:**
  - Linear: for each dimension
    - Find the highest low point, lowest high point; difference is separation
    - Normalize: divide separation by extent of that dimension
    - Across all dimensions, choose pair with biggest normalized separation
  - Quadratic: For each pair E1 and E2, calculate:
    - the rectangle  $J = \text{MBR}(E1, E2)$
    - $d = \text{area}(J) - (\text{area}(E1) + \text{area}(E2))$  (inefficiency)
    - Choose the pair with the largest d



## R-trees: Variations

- **There are many variations on R-trees**
  - Some change the key type (e.g. bounding spheres, “holey” bounding boxes)
  - Some fiddle with picksplit
  - Some complicate the structure (i.e. not exactly GiST)
  - And more...
- **Why so many?**
  - What’s wrong with R-trees
  - How good are R-trees anyway?



## GiST Performance

- **B+-trees have  $O(\log n)$  performance**
- **R-trees, RD-trees have no such guarantee**
  - search may have to traverse multiple paths
  - worst-case  $O(2n)$  to traverse entire tree
  - aggravated by random I/O
- **SO: when does it pay to build/use/invent an index? “Indexability”**
- **Basic questions:**
  1. Can data that’s co-retrieved be put together in leaf pages?
  2. Can an efficient “directory” be built on top?
- **Often, if (1) is possible, (2) rests on the key “shape” being accurate**



## The Gist of the GiST

- **Boil search trees down to their essence**
  - this is the “right” way to think about tree indexes
  - details of B+-trees, etc. are important, but not fundamental
  - the main idea is a hierarchy of clusters & labels, which grows by splitting bottom-up.
- **Unify B+-tree, R-tree, etc. in one ADT.**
  - code reuse!
- **Extensible in terms of data *and queries*.**
- **Raises nice theoretical questions of indexability.**



## More on GiST

- **Implemented in PostgreSQL**
  - Including high-concurrency and recovery
  - PostgreSQL include GiST extensions for R-trees over boxes/polygons/circles
    - Basis of the PostGIS Geographic Info System
  - Also includes RD-trees for text search
    - Recommended for indexing transactional text data
    - Use inverted (“GIN”) indexes for sloppy text data
- **Was implemented in Informix**
  - Purchased by IBM
- **More? <http://gist.cs.berkeley.edu>**