


Generalized Search Trees and Spatial Indexing

(Ramakrishnan 28.1, 28.2,
28.6)






Other Search Trees

- **Question: Can B+-trees handle more complicated searches?**
 - A typical example: "*gpa* > 3.7 *and* *age* < 18"
 - Same thing: "all restaurants in downtown Berkeley"
 - Even fancier: "all pictures resembling *</tmp/sunset.gif>*"
 - (Easy: "all pictures identical to *</tmp/sunset.gif>*")
- **B+-trees exploit data order to do range search**
 - 1-d range search is not always what you want





Search Trees in General

- **Lots of trees invented for multidimensional data**
 - e.g. R-trees, R*-trees, hB-trees, UB-trees, X-trees, etc. etc.
- **New tree indexes for other kinds of data**
 - Image/video search, timeseries matching, DNA sequence matching, etc.
- **Many are “unordered” B+-trees, fancy keys**
 - In some ways, B+-tree is just a “special case” of these trees.



Search Trees from 30,000 feet







A “Disorderly” B+-tree


- **How to search a disorderly B+-tree?**
 - Equality search is identical to traditional!
 - Range search = traversing multiple paths
 - follow all pointers where key range overlaps query range.

Traditional




Disorderly





Another Disorderly B+-Tree



- **Insert 41???**
 - keys need not cover all possible values
- **Search for 77??**
 - keys may “overlap”



Generalized Search Tree (GiST)

- **A disorderly B+-tree, with *user-defined* keys**
 - tree doesn't interpret keys.
 - “user” implements keys as an OO class, with methods that guide search, insert, delete, split, etc.
- **Structure: balanced tree of (p , ptr) pairs**
 - p is an index key or “predicate”
 - p holds for all data records below ptr
 - need n keys for n pointers (unlike B+-tree)




User-provided Key Methods

- **Search:**
 - **Consistent**(E, q): $E.p \wedge q$? (no/maybe)
- **Generating new keys after splits:**
 - **Union**(P): new key that holds for all tuples in P
- **Data organization:**
 - **Penalty**(E_1, E_2):
“badness” of inserting E_2 in subtree E_1
 - **PickSplit**(P): split P into two groups of entries



Search

- **General technique:**
 - Depth-First Search where **Consistent** is TRUE
- **Incremental algorithm:**
 - Maintain a search stack of `<page_id, offset>` pairs from root to current data entry
 - Each “get next” call moves the offset to right
 - When nothing left on page, pop stack, move right in parent (or recurse) and go down to leaf.




Insert

- **descend tree along least increase in Penalty**
- **if there's room at leaf, insert there**
- **else split according to PickSplit**
 - do B-tree-style recursive splitting
- **propagate changes (recursively) using Union**
 - make sure all ancestor keys are consistent with inserted item


Q: what happened to B-tree “copy up” and “push up” logic?

Let's revisit with example of R-tree in a few slides

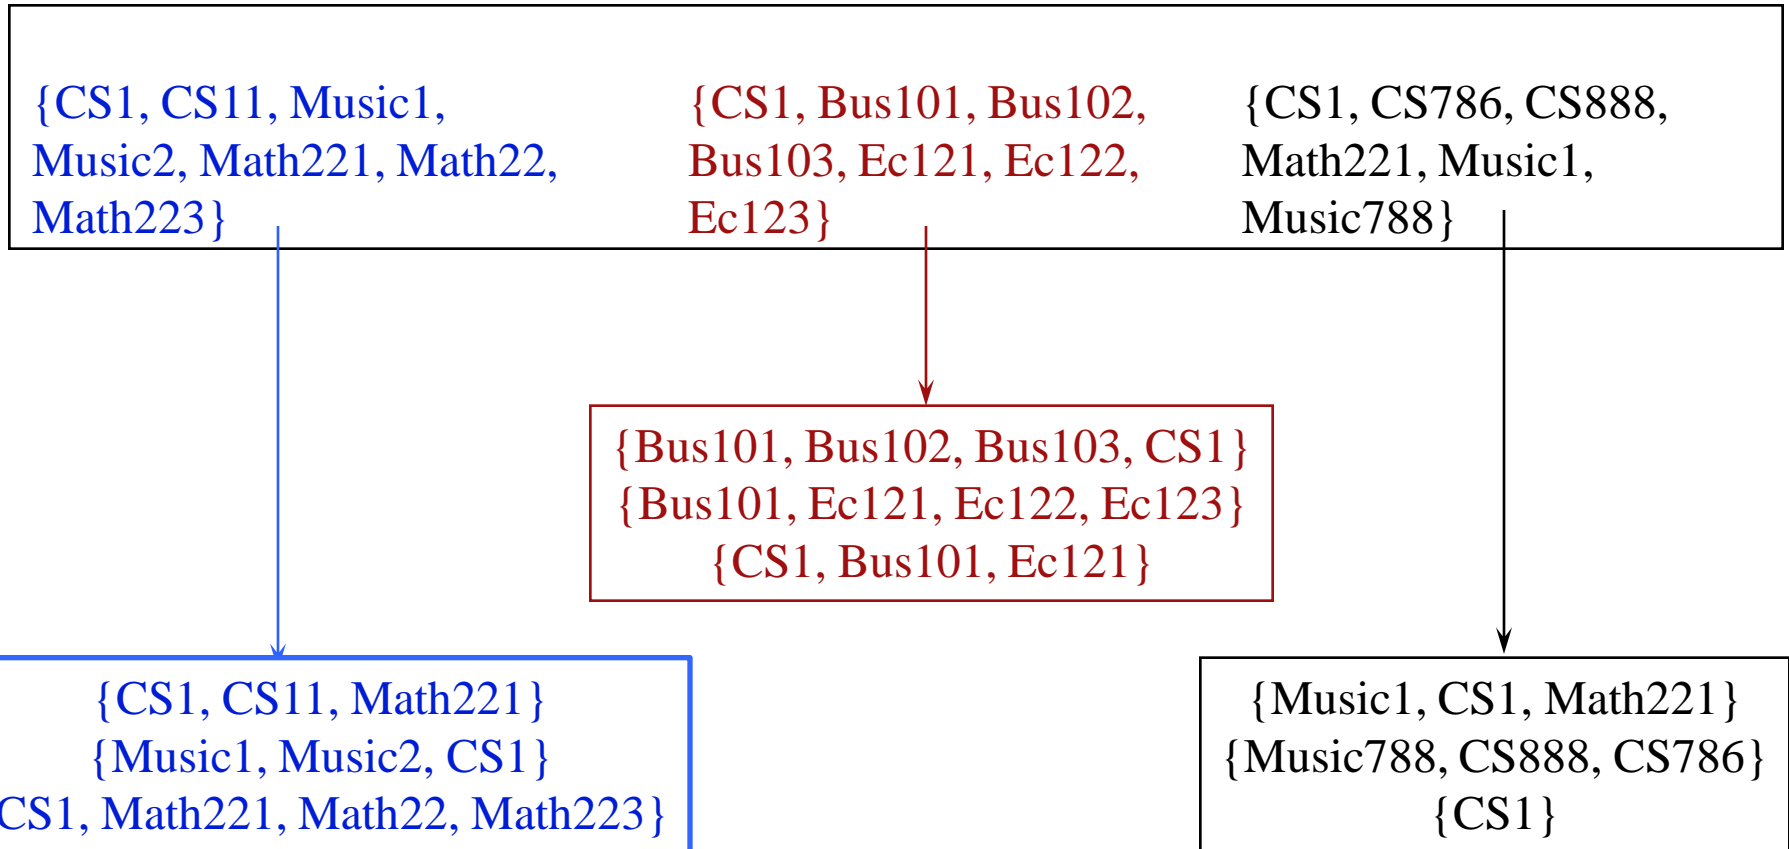



Delete

- **find the entry via Search, and delete it**
- **propagate changes (recursively) using Union**
- **on underflow:**
 - reinsert stuff on page and delete page
 - why not borrow/merge a la B+-trees?



RD-tree: GiST for Sets






RD-trees

- **Logically, keys represent minimal supersets**
- **Queries: Contains, Intersects, Equals...**
- **Consistent(E, q):**
 - Varies slightly on query type ... e.g. $E.p \cap q \neq \emptyset$
- **Union(P): set-union of keys**
- **Penalty(E, F): $|E.p \cup F.p| - |E.p|$**
- **PickSplit(P): many possible algorithms**
 - One goal: minimize sum of cardinalities of 2 pages

Used in Postgres as an alternative to inverted indexes

Note: need key compression here (like postings lists)




R-tree: A GiST over 2-d data


- **(Invented ~10 years before GiST or RD-Tree)**
 - Also at Berkeley ;-)
- **Logically, keys represent “bounding boxes”**
- **Queries: Contains, Overlaps, Equals ... *bbox***
- **Consistent(E, q):**
 - Varies slightly on query type: e.g. $E.p$ overlaps q ?
- **Union(P): bounding box of all entries**
- **Penalty(E, F): $\text{size}(\text{Union}(\{E, F\})) - \text{size}(E)$**
- **PickSplit(P):**
 - goal: minimize sum of areas of the 2 pages



R-trees, Slowly. Problem:


- **Given a collection of geometric objects (points, lines, polygons, ...)**
- **organize them on disk, to answer spatial queries (range, near neighbors, etc)**






R-trees


- **Main idea: a 2-D B-tree with overlapping keys!**
 - => guaranteed 50% utilization
 - => easier insertion/split algorithms.
 - (only deal with Minimum Bounding Rectangles - **MBRs**)






R-tree Structure


- **A multi-way tree of disk blocks**
- **Index nodes and data (leaf) nodes**
- **All leaf nodes appear on the same level**
- **Every node contains between m and M entries**
- **The root node has at least 2 entries (children)**



Example


- **eg., w/ fanout 4: group nearby rectangles to parent MBRs; each group -> disk page**






Example


- **F=4**






Example


- **F=4**






R-trees - format of nodes


- **{(MBR; obj_ptr)}** for leaf nodes







R-trees - format of nodes


- **{(MBR; node_ptr)}** for non-leaf nodes



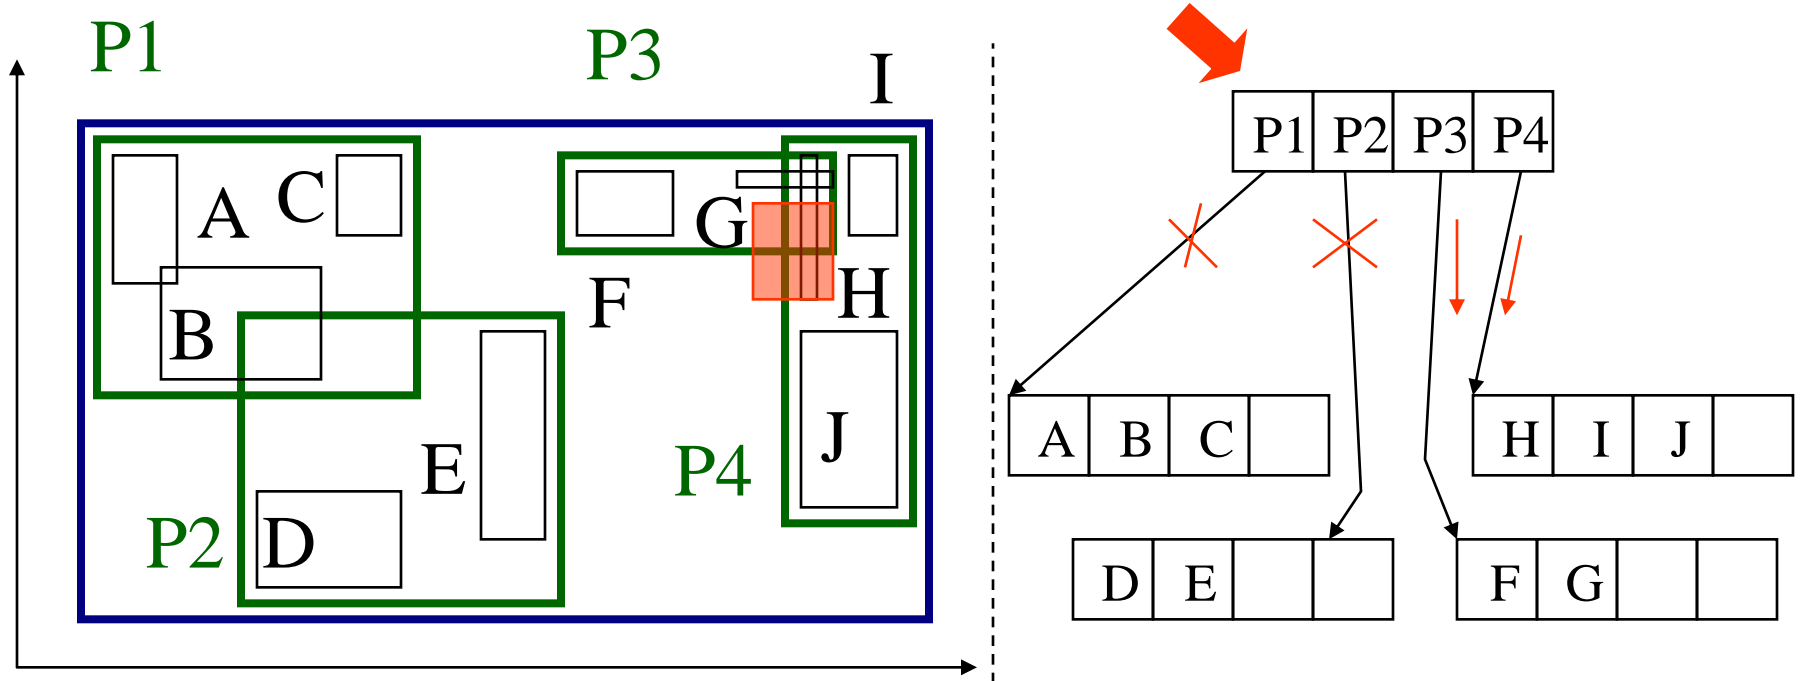



R-trees: Search





R-trees: Search






R-trees:Search


- **Main points:**


- every parent node completely covers its 'children'
- a child MBR may be covered by more than one parent - it is stored under ONLY ONE of them. (ie., no need for dup. elim.)
- a point query may follow multiple branches.
- everything works for **any(?)** dimensionality



R-trees: Near Neighbor Search


- **Rather than a stack (for depth-first), maintain a priority queue of nodes to visit**
- **Upon visiting a page, load all the children to be traversed into the priority queue**
 - Priority = MinDistance






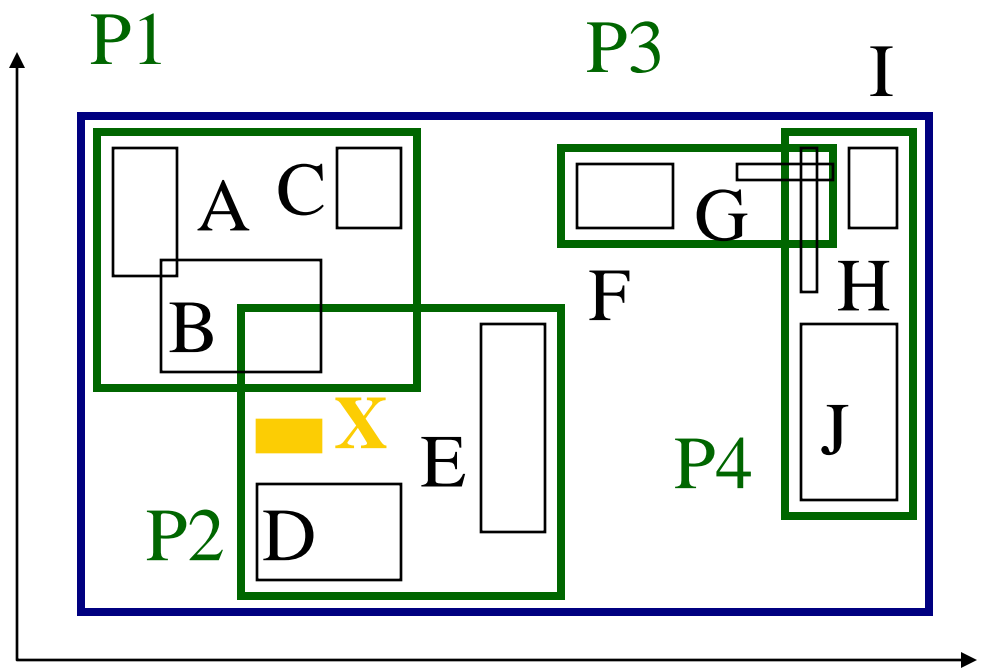
R-tree: Incremental Near Neighbors


```
PQ.insert(root)
while (!PQ.empty)
  next = PQ.deletemin
  if (next is a data entry) return (next)
  else // next is a page in the index
    for each item i on next
      PQ.insert(i)
    end for
  end if
end while
```



R-trees: Insertion



Insert X






R-trees: Insertion


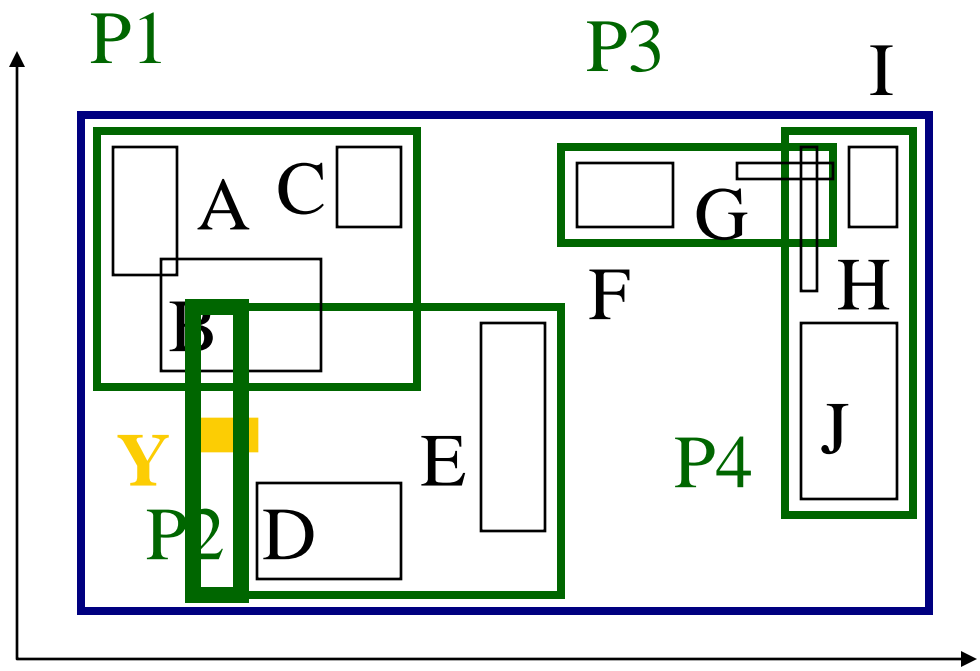
Insert Y






R-trees: Insertion


- **Extend the parent MBR**






R-trees: Insertion


- **How to find the next node to insert the new object?**
 - Penalty metric: At each level, find the entry that needs the least enlargement to include Y. Resolve ties using the area (smallest)
- **Other methods have been proposed**
 - E.g. perhaps useful to minimize the perimeter of MBRs too?



R-trees: Insertion


- If node is full then **Split** : ex. Insert w






R-trees: PickSplit

- **Split node P1: partition the MBRs into two groups.**





- Naïve: exhaustive
 - cost?
- quadratic split
- linear split



R-tree PickSplit: idea


- pick two rectangles as 'seeds';
- assign each rectangle 'R' to the 'closest' 'seed'





R-trees: Split

- **pick two rectangles as 'seeds';**
- **assign each rectangle 'R' to the 'closest' 'seed':**
- **'closest': the smallest increase in area**





R-trees: Split

- **How to pick Seeds:**

- Linear: for each dimension

- Find the highest low point, lowest high point; difference is separation
 - Normalize: divide separation by extent of that dimension
 - Across all dimensions, choose pair with biggest normalized separation

- Quadratic: For each pair E1 and E2, calculate:

- the rectangle $J = \text{MBR}(E1, E2)$
 - $d = \text{area}(J) - (\text{area}(E1) + \text{area}(E2))$ (inefficiency)
 - Choose the pair with the largest d



R-trees: Variations

- **There are many variations on R-trees**
 - Some change the key type (e.g. bounding spheres, “holey” bounding boxes)
 - Some fiddle with picksplit
 - Some complicate the structure (i.e. not exactly GiST)
 - And more...
- **Why so many?**
 - What’s wrong with R-trees
 - How good are R-trees anyway?



GiST Performance

- **B+-trees have $O(\log n)$ performance**
- **R-trees, RD-trees have no such guarantee**
 - search may have to traverse multiple paths
 - worst-case $O(2n)$ to traverse entire tree
 - aggravated by random I/O
- **SO: when does it pay to build/use/invent an index?**
“Indexability”
- **Basic questions:**
 1. Can data that’s co-retrieved be put together in leaf pages?
 2. Can an efficient “directory” be built on top?
- **Often, if (1) is possible, (2) rests on the key “shape” being accurate**



The Gist of the GiST

- **Boil search trees down to their essence**
 - this is the “right” way to think about tree indexes
 - details of B+-trees, etc. are important, but not fundamental
 - the main idea is a hierarchy of clusters & labels, which grows by splitting bottom-up.
- **Unify B+-tree, R-tree, etc. in one ADT.**
 - code reuse!
- **Extensible in terms of data *and queries*.**
- **Raises nice theoretical questions of indexability.**



More on GiST

- **Implemented in PostgreSQL**
 - Including high-concurrency and recovery
 - PostgreSQL include GiST extensions for R-trees over boxes/polygons/circles
 - Basis of the PostGIS Geographic Info System
 - Also includes RD-trees for text search
 - Recommended for indexing transactional text data
 - Use inverted ("GIN") indexes for sloppy text data
- **Was implemented in Informix**
 - Purchased by IBM
- **More? <http://gist.cs.berkeley.edu>**