


## Relational Algebra

R & G, Chapter 4

By relieving the brain of all unnecessary work, a good notation sets it free to concentrate on more advanced problems, and, in effect, *increases the mental power of the race.*

-- Alfred North Whitehead (1861 - 1947)



## Relational Query Languages

- **Query languages:** Allow manipulation and **retrieval of data** from a database.
- **Relational model supports simple, powerful QLS:**
  - Strong formal foundation based on logic.
  - Allows for much optimization.
- **Query Languages != programming languages!**
  - QLS not expected to be "Turing complete".
  - QLS not intended to be used for complex calculations.
  - QLS support easy, efficient access to large data sets.



## Formal Relational Query Languages

Two mathematical Query Languages form the basis for "real" languages (e.g. SQL), and for implementation:

**Relational Algebra:** More **operational**, very useful for representing execution plans.

**Relational Calculus:** Lets users describe what they want, rather than how to compute it. (**Non-procedural, declarative.**)

☒ *Understanding Algebra & Calculus is key to understanding SQL, query processing!*



## Preliminaries

- A query is applied to **relation instances**, and the result of a query is also a **relation instance**.
  - **Schemas** of input relations for a query are **fixed** (but query will run over any legal instance)
  - The **schema for the result** of a given query is also **fixed**. It is determined by the definitions of the query language constructs.
- **Positional vs. named-field notation:**
  - Positional notation easier for formal definitions, named-field notation more readable.
  - Both used in SQL
    - Though positional notation is not encouraged



## Relational Algebra: 5 Basic Operations

- **Selection** ( $\sigma$ ) Selects a subset of **rows** from relation (horizontal).
- **Projection** ( $\pi$ ) Retains only wanted **columns** from relation (vertical).
- **Cross-product** ( $\times$ ) Allows us to combine two relations.
- **Set-difference** ( $-$ ) Tuples in r1, but not in r2.
- **Union** ( $\cup$ ) Tuples in r1 or in r2.

Since each operation returns a relation, **operations can be composed!** (Algebra is "closed".)



## Example Instances

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Boats

<u>bid</u>	<u>bname</u>	<u>color</u>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

S1

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



## Projection ( $\pi$ )

- Examples:  $\pi_{age}(S2)$  ;  $\pi_{sname, rating}(S2)$
- Retains only attributes that are in the "projection list".
- Schema** of result:
  - exactly the fields in the projection list, with the same names that they had in the input relation.
- Projection operator has to **eliminate duplicates** (How do they arise? Why remove them?)
  - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it. (Why not?)



## Projection ( $\pi$ )

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

 $\pi_{sname, rating}(S2)$ 

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

age
35.0
55.5

 $\pi_{age}(S2)$ 


## Selection ( $\sigma$ )

- Selects rows that satisfy **selection condition**.
- Result is a relation.
- Schema** of result is same as that of the input relation.
- Do we need to do duplicate elimination?

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

 $\sigma_{rating > 8}(S2)$ 

sname	rating
yuppy	9
rusty	10

 $\pi_{sname, rating}(\sigma_{rating > 8}(S2))$ 


## Union and Set-Difference

- Both of these operations take two input relations, which must be **union-compatible**:
  - Same number of fields.
  - Corresponding fields have the same type.
- For which, if any, is duplicate elimination required?



## Union

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

 $S1 \cup S2$ 

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2



## Set Difference

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

sid	sname	rating	age
22	dustin	7	45.0

 $S1 - S2$ 

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

sid	sname	rating	age
28	yuppy	9	35.0
44	guppy	5	35.0

 $S2 - S1$



## Cross-Product

- $S1 \times R1$ : Each row of S1 paired with each row of R1.
- Q: How many rows in the result?
- **Result schema** has one field per field of S1 and R1, with field names 'inherited' if possible.
  - *May have a naming conflict*: Both S1 and R1 have a field with the same name.
  - In this case, can use the *renaming operator*:

$$\rho(C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$$



## Cross Product Example

sid	bid	day
22	101	10/10/96
58	103	11/12/96

R1

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

$S1 \times R1 =$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96



## Compound Operator: Intersection

- In addition to the 5 basic operators, there are several additional "Compound Operators"
  - These add no computational power to the language, but are useful shorthands.
  - Can be expressed solely with the basic ops.
- Intersection takes two input relations, which must be **union-compatible**.
- Q: How to express it using basic operators?

$$R \cap S = R - (R - S)$$



## Intersection

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$$S1 \cap S2$$



## Compound Operator: Join

- Joins are compound operators involving cross product, selection, and (sometimes) projection.
- Most common type of join is a "**natural join**" (often just called "join").  $R \bowtie S$  conceptually is:
  - Compute  $R \times S$
  - Select rows where attributes that appear in both relations have equal values
  - Project all unique attributes and one copy of each of the common ones.
- Note: Usually done much more efficiently than this.



## Natural Join Example

sid	bid	day
22	101	10/10/96
58	103	11/12/96

R1

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

$S1 \bowtie R1 =$

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96



## Other Types of Joins

- Condition Join (or "theta-join"):

$$R \bowtie_c S = \sigma_c(R \times S)$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

- **Result schema** same as that of cross-product.
- May have fewer tuples than cross-product.
- **Equi-Join**: Special case: condition *c* contains only conjunction of equalities.



## Examples

Reserves

sid	bid	day
22	101	10/10/96
58	103	11/12/96

Sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Boats

bid	bname	color
101	Interlake	Blue
102	Interlake	Red
103	Clipper	Green
104	Marine	Red



Find names of sailors who've reserved boat #103

- **Solution 1:**  $\pi_{sname}((\sigma_{bid=103} Reserves) \bowtie Sailors)$
- **Solution 2:**  $\pi_{sname}(\sigma_{bid=103}(Reserves \bowtie Sailors))$



Find names of sailors who've reserved a red boat

- **Information about boat color only available in Boats; so need an extra join:**

$$\pi_{sname}((\sigma_{color=red} Boats) \bowtie Reserves \bowtie Sailors)$$

❖ A more efficient solution:

$$\pi_{sname}(\pi_{sid}(\pi_{bid}(\sigma_{color=red} Boats) \bowtie Res) \bowtie Sailors)$$

☒ A query optimizer can find this given the first solution!



Find sailors who've reserved a red or a green boat

- **Can identify all red or green boats, then find sailors who've reserved one of these boats:**

$$\rho(Tempboats, (\sigma_{color=red \vee color=green} Boats))$$


$$\pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$$



Find sailors who've reserved a red and a green boat

- **Cut-and-paste previous slide?**

$$\pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$$





Find sailors who've reserved a red and a green boat

- **Previous approach won't work! Must identify sailors who've reserved red boats, sailors who've reserved green boats, then find the intersection (note that *sid* is a key for Sailors):**

$$\rho(Tempred, \pi_{sid}((\sigma_{color='red'} Boats) \bowtie Reserves))$$

$$\rho(Tempgreen, \pi_{sid}((\sigma_{color='green'} Boats) \bowtie Reserves))$$

$$\pi_{sname}((Tempred \cap Tempgreen) \bowtie Sailors)$$


## Summary

- **Relational Algebra: a small set of operators mapping relations to relations**
  - Operational, in the sense that you specify the explicit order of operations
  - A *closed* set of operators! Can mix and match.
- **Basic ops include:**  $\sigma, \pi, \times, \cup, -$
- **Important compound ops:**  $\cap, \bowtie$