

Implementation of Relational Operations

R&G - Chapters 12 and 14




Introduction

- Today's topic: **QUERY PROCESSING**
- Some database operations are **EXPENSIVE**
- Can greatly improve performance by being "smart"
 - e.g., can speed up 1,000,000x over naïve approach
- Main weapons are:
 1. clever implementation techniques for operators
 2. exploiting relational algebra "equivalences"
 3. using statistics and cost models to choose among these.

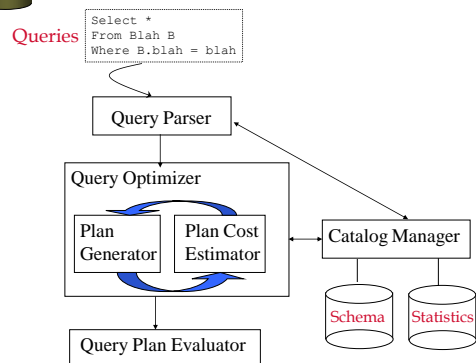


A Really Bad Query Optimizer

- For each **Select-From-Where** query block
 - Create a plan that:
 - Forms the cross product of the FROM clause
 - Applies the WHERE clause
- Then, as needed:
 - Apply the GROUP BY clause
 - Apply the HAVING clause
 - Apply any projections and output expressions
 - Apply duplicate elimination and/or ORDER BY



Cost-based Query Sub-System



The Query Optimization Game

- Goal is to pick a "good" plan
 - Good = low expected cost, under *cost model*
 - Degrees of freedom:
 - access methods
 - physical operators
 - operator orders
- Roadmap for this topic:
 - First: implementing individual operators
 - Then: optimizing multiple operators



Relational Operations

- We will consider how to implement:
 - Selection (σ) Select a subset of rows.
 - Projection (π) Remove unwanted columns.
 - Join (\bowtie) Combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 and in reln. 2.
- Q: What about Intersection?



Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- **Similar to old schema; *rname* added for variations.**
- **Sailors:**
 - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
 - $|S|=500$, $p_S=80$.
- **Reserves:**
 - Each tuple is 40 bytes, 100 tuples per page, 1000 pages.
 - $|R|=1000$, $p_R=100$.



Our options ...

- **If no appropriate index exists:**
Must scan the whole relation
cost = $|R|$. For "reserves" = 1000 I/Os.



Simple Selections

$\sigma_{R.attr \text{ op } value}(R)$


```
SELECT *
FROM   Reserves R
WHERE  R.rname < 'C%'
```

- **How best to perform? Depends on:**
 - what indexes are available
 - expected size of result
- **Size of result approximated as**
(size of *R*) * selectivity
 - selectivity estimated via statistics – we will discuss shortly.




Our options ...

- **With index on selection attribute:**
 1. Use index to find qualifying data entries
 2. Retrieve corresponding data records
- Total cost = cost of step 1 + cost of step 2**
- For "reserves", if selectivity = 10% (100 pages, 10000 tuples):
 - If **clustered** index, cost is a little over **100** I/Os;
 - If **unclustered**, could be up to **10000** I/Os! ... unless ...



Refinement for unclustered indexes

1. Find qualifying data entries.
 2. Sort the rid's of the data records to be retrieved.
 3. Fetch rids in order.
- Each data page is looked at just once (though # of such pages likely to be higher than with clustering).



General Selection Conditions

✉ (*day*<8/9/94 AND *rname*='Paul') OR *bid*=5 OR *sid*=3

- **First, convert to conjunctive normal form (CNF):**
 - (*day*<8/9/94 OR *bid*=5 OR *sid*=3) AND (*rname*='Paul' OR *bid*=5 OR *sid*=3)
- **We only discuss the case with no ORs**
- **Terminology:**
 - A **B-tree** index **matches** terms that involve only attributes in a **prefix** of the search key. e.g.:
 - Index on **<a, b, c>** matches **a=5 AND b=3**, but not **b=3**.



2 Approaches to General Selections

Approach I:

1. Find the *cheapest access path*
2. retrieve tuples using it
3. Apply any remaining terms that don't **match** the index

- *Cheapest access path*: An index or file scan that we estimate will require the fewest page I/Os.



Cheapest Access Path - Example

query: *day < 8/9/94 AND bid=5 AND sid=3*

some options:

*B+tree index on day; check bid=5 and sid=3 afterward.
hash index on <bid, sid>; check day<8/9/94 afterward.*

- How about a B+tree on <rname,day>?
- How about a B+tree on <day, rname>?
- How about a Hash index on <day, rname>?



2 Approaches to General Selections

Approach II: use 2 or more matching indexes.

1. From each index, get set of rids
2. Compute intersection of rid sets
3. Retrieve records for rids in intersection
4. Apply any remaining terms

EXAMPLE: *day<8/9/94 AND bid=5 AND sid=3*

Suppose we have an index on *day*, and another index on *sid*.

- Get rids of records satisfying *day<8/9/94*.
- Also get rids of records satisfying *sid=3*.
- Find intersection, then retrieve records, then check *bid=5*.



Projection

SELECT	DISTINCT
	R.sid, R.bid
FROM	Reserves R

- Issue is removing **duplicates**.
- Use **sorting!!**
 1. Scan R, extract only the needed attributes
 2. Sort the resulting set
 3. Remove adjacent duplicates

Cost:

Reserves with size ratio 0.25 = 250 pages.

With 20 buffer pages can sort in 2 passes, so:

$$1000 + 250 + 2 * 2 * 250 + 250 = 2500 \text{ I/Os}$$



Projection -- improved

- **Modify the external sort algorithm:**
 - Modify Pass 0 to eliminate unwanted fields.
 - Modify Passes 1+ to eliminate duplicates.

Cost:

Reserves with size ratio 0.25 = 250 pages.

With 20 buffer pages can sort in 2 passes, so:

1. Read 1000 pages
2. Write 250 (in runs of 40 pages each)
3. Read and merge runs

$$\text{Total cost} = 1000 + 250 + 250 = 1500.$$



Other Projection Tricks

If an index search key contains all wanted attrs:

- Do **index-only scan**
 - Apply projection techniques to data entries (*much smaller!*)

If a B+Tree index search key prefix has all wanted attrs:


- Do **in-order index-only scan**
 - Compare adjacent tuples on the fly (*no sorting required!*)



Query Execution Framework

```
SELECT DISTINCT name, gpa
FROM Students
```

One possible **query execution plan**:



Example: Sort

```
class Sort extends iterator {
    void init();
    tuple next();
    void close();
    iterator &inputs[1];
    int numberOfRuns;
    DiskBlock runs[];
    RID nextRID[];
}
```

- **init():**
 - generate the sorted runs on disk (*passes 0 to n-1*)
 - Allocate runs[] array and fill in with disk pointers.
 - Initialize numberOfRuns
 - Allocate nextRID array and initialize to first RID of each run
- **next():**
 - nextRID array tells us where we're "up to" in each run
 - find the next tuple to return based on nextRID array
 - advance the corresponding nextRID entry
 - return tuple (or EOF -- "End of Run" -- if no tuples remain)
- **close():**
 - deallocate the runs and nextRID arrays



Iterators

Iterator

- **Relational operators are all subclasses of the class iterator:**

```
class iterator {
    void init();
    tuple next();
    void close();
    iterator &inputs[];
    // additional state goes here
}
```

- **Note:**
 - Edges in the graph are specified by inputs (max 2, usually)
 - Any iterator can be input to any other!



Postgres Version

- **src/backend/executor/nodeSort.c**
 - ExecInitSort (init)
 - ExecSort (next)
 - ExecEndSort (close)
- **The encapsulation stuff is hardwired into the Postgres C code**
 - Postgres predates even C++!
 - See src/backend/execProcNode.c for the code that "dispatches the methods" explicitly!



Joins

```
SELECT *
FROM Reserves R1, Sailors S1
WHERE R1.sid=S1.sid
```

- **Joins are very common.**
- **$R \times S$ is large; so, $R \times S$ followed by a selection is inefficient.**
- **Many approaches to reduce join cost.**
- **Join techniques we will cover today:**
 1. Nested-loops join
 2. Index-nested loops join
 3. Sort-merge join



Simple Nested Loops Join

```
R ⋈ S: foreach tuple r in R do
    foreach tuple s in S do
        if r1 == s1 then add <r, s> to result
```

Cost = $(p_R * |R|) * |S| + |R| = 100 * 1000 * 500 + 1000$ IOs
 – At 10ms/IO, Total time: ???

- **What if smaller relation (S) was "outer"?**
- **What assumptions are being made here?**
- **What is cost if one relation can fit entirely in memory?**

Page-Oriented Nested Loops Join


```
R ⋈ S: foreach page  $b_R$  in R do
    foreach page  $b_S$  in S do
        foreach tuple  $r$  in  $b_R$  do
            foreach tuple  $s$  in  $b_S$  do
                if  $r_1 == s_1$  then add  $\langle r, s \rangle$  to result
```

Cost = $|R| * |S| + |R| = 1000 * 500 + 1000$

- If smaller relation (S) is outer, cost = $500 * 1000 + 500$
- Much better than naïve per-tuple approach!

Block Nested Loops Join

- Page-oriented NL doesn't exploit extra buffers :(
- Idea to use memory efficiently:

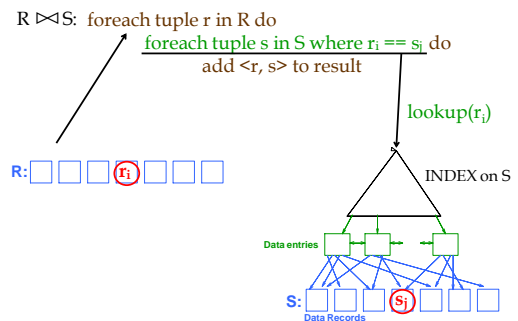


Cost: Scan outer + (#outer blocks * scan inner)
 #outer blocks = $\lceil \# \text{ of pages of outer} / \text{blocksize} \rceil$

Examples of Block Nested Loops Join

- Say we have **B = 100+2** memory buffers
- **Join cost = |outer| + (#outer blocks * |inner|)**
 #outer blocks = $|outer| / 100$
- **With R as outer (|R| = 1000):**
 - Scanning R costs 1000 IO's (done in 10 blocks)
 - Per block of R, we scan S; costs $10 * 500$ IO's
 - Total = $1000 + 10 * 500$.
- **With S as outer (|S| = 500):**
 - Scanning S costs 500 IO's (done in 5 blocks)
 - Per block of S, we scan R; costs $5 * 1000$ IO's
 - Total = $500 + 5 * 1000$.

Index Nested Loops Join



Index Nested Loops Join

```
R ⋈ S: foreach tuple r in R do
    foreach tuple s in S where  $r_1 == s_1$  do
        add  $\langle r, s \rangle$  to result
```

Cost = $|R| + (|R| * p_R) * \text{cost to find matching S tuples}$

- If index uses Alt. 1, cost = cost to traverse tree from root to leaf.
- For Alt. 2 or 3:
 1. Cost to lookup RID(s); typically 2-4 IO's for B+Tree.
 2. Cost to retrieve records from RID(s); depends on clustering.
 - Clustered index: 1 I/O per page of matching S tuples.
 - Unclustered: up to 1 I/O per matching S tuple.

Sort-Merge Join

1. Sort R on join attr(s)
2. Sort S on join attr(s)
3. Scan sorted-R and sorted-S in tandem, to find matches

Example:

```
SELECT *
FROM Reserves R1, Sailors S1
WHERE R1.sid=S1.sid
```

sid	bid	day	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin



Cost of Sort-Merge Join

- **Cost: Sort R + Sort S + ($|R| + |S|$)**
 - But in worst case, last term could be $|R| * |S|$ (*very unlikely!*)
 - Q: what is worst case?

Suppose **B = 35** buffer pages:

- Both R and S can be sorted in 2 passes
- Total join cost = $4 * 1000 + 4 * 500 + (1000 + 500) = 7500$

Suppose **B = 300** buffer pages:

- Again, both R and S sorted in 2 passes
- Total join cost = **7500**

Block-Nested-Loop cost = 2500 ... 15,000



Other Considerations ...

1. An important refinement:

Do the join during the final merging pass of sort!

- If have enough memory, can do:
 1. Read R and write out sorted runs
 2. Read S and write out sorted runs
 3. Merge R-runs and S-runs, and find R \bowtie S matches

Cost = $3 * |R| + 3 * |S|$

Q: how much memory is "enough" (*will answer next time ...*)

2. Sort-merge join an especially good choice if:

- one or both inputs are **already sorted** on join attribute(s)
- output is **required to be sorted** on join attributes(s)

Q: how to take these savings into account? (*stay tuned ...*)



Summary

- A virtue of relational DBMSs:

queries are composed of a few basic operators

- The implementation of these operators can be **carefully tuned**
- Many alternative implementation techniques for each operator
 - No universally superior technique for most operators.
- Must consider available alternatives
 - Called "Query optimization" -- we will study this topic soon!