


# Implementation of Relational Operations (Part 2)

R&G - Chapters 12 and 14



## An Alternative to Sorting: Hashing!


- **Idea:**
  - Many of the things we use sort for don't exploit the *order* of the sorted data
  - e.g.: removing duplicates in DISTINCT
  - e.g.: finding matches in JOIN
- **Often good enough to match all tuples with equal values**
- **Hashing does this!**
  - And may be cheaper than sorting! (Hmmm....!)
  - But how to do it for data sets bigger than memory??

## General Idea


- **Two phases:**
  - **Partition:** use a hash function  $h$  to split tuples into partitions on disk.
    - Key property: all matches live in the same partition.
  - **ReHash:** for each partition on disk, build a main-memory hash table using a hash function  $h_2$

## Two Phases

### • Partition:




### • Rehash:




## Duplicate Elimination using Hashing

- read one bucket at a time
- for each group of identical tuples, output one



## Cost of External Hashing



$$\text{cost} = 4 * [R] \text{ IO's}$$



## Memory Requirement


- **How big of a table can we hash in two passes?**
  - B-1 "partitions" result from Phase 0
  - Each should be no more than B pages in size
  - Answer:  $B(B-1)$ .
  - Said differently:*  
We can hash a table of size N pages in about  $\sqrt{N}$  space
  - *Note: assumes hash function distributes records evenly!*
- **Have a bigger table? Recursive partitioning!**



How does this compare with  
external sorting?




## Cost of External Hashing



cost =  $4 \cdot [R]$  IO's



## Cost of External Sorting



cost =  $4 \cdot [R]$  IO's



## Memory Requirement for External Sorting

- **How big of a table can we sort in two passes?**
  - Each "sorted run" after Phase 0 is of size B
  - Can merge up to B-1 sorted runs in Phase 1
  - Answer:  $B(B-1)$ .
  - Said differently:*  
We can sort a table of size N pages in about  $\sqrt{N}$  space
- **Have a bigger table? Additional merge passes!**



## So which is better ??

- **Based on our simple analysis:**
  - Same memory requirement for 2 passes
  - Same IO cost
- **Digging deeper ...**
- **Sorting pros:**
  - Great if input already sorted (or *almost* sorted)
  - Great if need output to be sorted anyway
  - Not sensitive to "data skew" or "bad" hash functions
- **Hashing pros:**
  - Highly parallelizable (will discuss later in semester)
    - So is sorting, with some work
  - Can exploit extra memory to reduce # IOs (*stay tuned...*)

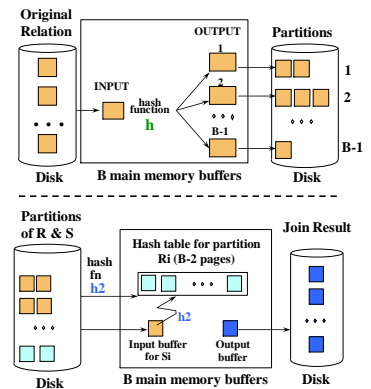


before we optimize hashing further ...

Q: Can we use hashing for JOIN ?



## Hash Join



## Cost of Hash Join

- **Partitioning phase:** read+write both relations  
 $\Rightarrow 2([R]+[S])$  I/Os
- **Matching phase:** read both relations, write output  
 $\Rightarrow [R]+[S] + [\text{output}]$  I/Os
- **Total cost of 2-pass hash join =  $3([R]+[S])+[\text{output}]$**

Q: what is cost of 2-pass *sort join*?


Q: how much memory needed for 2-pass *sort join*?

Q: how much memory needed for 2-pass *hash join*?



## An important optimization to hashing

- Have B memory buffers
- Want to hash relation of size N




If  $B < N < B^2$ , will have unused memory ...



## Hybrid Hashing

- **Idea:** keep one of the hash buckets in memory!




Q: how do we choose the value of *k*?



## Cost reduction due to hybrid hashing

- **Now:**





## Summary: Hashing vs. Sorting

- **Sorting pros:**
  - Good if input already sorted, or need output sorted
  - Not sensitive to *data skew* or *bad hash functions*
- **Hashing pros:**
  - Often cheaper due to *hybrid hashing*
  - For join: # passes depends on size of *smaller* relation
  - Highly *parallelizable*