


Relational Query Optimization

CS186

R & G Chapters 12/15




Review

- **Implementation of single Relational Operations**
- **Choices depend on indexes, memory, stats,...**
- **Joins**
 - Blocked nested loops:
 - simple, exploits extra memory
 - Indexed nested loops:
 - best if 1 rel small and one indexed
 - Sort/Merge Join
 - good with small amount of memory, bad with duplicates
 - Hash Join
 - fast (enough memory), bad with skewed data



Query Optimization Overview

- Query can be converted to relational algebra
- Rel. Algebra converted to tree, joins as branches
- **Each operator has implementation choices**
- **Operators can also be applied in different order!**

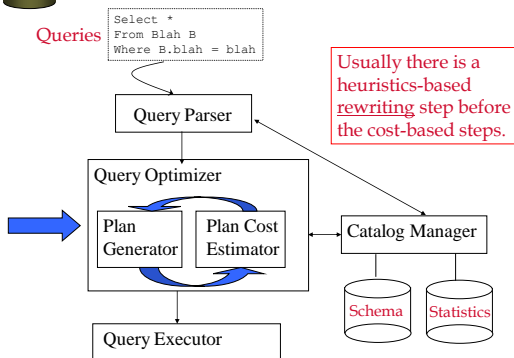


Query Optimization Overview (cont.)

- **Plan:** *Tree of R.A. ops (and some others) with choice of algorithm for each op.*
 - Recall: Iterator interface (next()!)
- **Three main issues:**
 - For a given query, **what plans** are considered?
 - How is the **cost of a plan** estimated?
 - How do we “search” in the “plan space”?
- **Ideally: Want to find best plan.**
- **Reality: Avoid worst plans!**



Cost-based Query Sub-System



Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
 Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)


- **As seen in previous lectures...**
- **Reserves:**
 - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
 - Assume there are 100 boats
- **Sailors:**
 - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
 - Assume there are 10 different ratings
- **Assume we have 5 pages in our buffer pool!**




Motivating Example

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```


- **Cost: 500+500*1000 I/Os**
- **By no means the worst plan!**
- **Misses several opportunities:** selections could have been 'pushed' earlier, no use is made of any available indexes, etc.
- **Goal of optimization:** To find more efficient plans that compute the same answer.




Alternative Plans – Push Selects (No Indexes)




Alternative Plans – Push Selects (No Indexes)



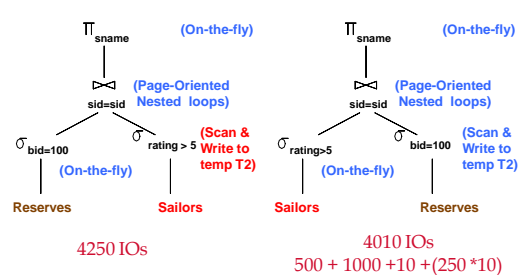
Alternative Plans – Push Selects (No Indexes)



Alternative Plans – Push Selects (No Indexes)




Alternative Plans – Push Selects (No Indexes)





More Alternative Plans (No Indexes)




- **Main difference: Sort Merge Join**
- **With 5 buffers, cost of plan:**
 - Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution) = 1010.
 - Scan Sailors (500) + write temp T2 (250 pages, if have 10 ratings) = 750.
 - Sort T1 ($2 \times 2 \times 10$) + sort T2 ($2 \times 4 \times 250$) + merge ($10 + 250$) = 2300
 - Total: 4060 page I/Os.
- **If use BNL join, join = $10 + 4 \times 250$, total cost = 2770.**
- **Can also 'push' projections, but must be careful!**
 - T1 has only *sid*, T2 only *sid*, *sname*:
 - T1 fits in 3 pgs, cost of BNL under 250 pgs, total < 2000.



More Alt Plans: Indexes

- **With clustered index on *bid* of Reserves, we get $100,000/100 = 1000$ tuples on $1000/100 = 10$ pages.**
- **INL with outer not materialized.**
 - Projecting out unnecessary fields from outer doesn't help.
- ❖ Join column *sid* is a key for Sailors.
 - At most one matching tuple, unclustered index on *sid* OK.
- ❖ Decision not to push *rating* > 5 before the join is based on availability of *sid* index on Sailors.
- ❖ **Cost:** Selection of Reserves tuples (10 I/Os); then, for each, must get matching Sailors tuple (1000×1.2); total 1210 I/Os.



What is needed for optimization?

- **A closed set of operators**
 - Relational ops (table in, table out)
 - Encapsulation based on iterators
- **Plan space, based on**
 - Based on relational equivalences, different implementations
- **Cost Estimation, based on**
 - Cost formulas
 - Size estimation, based on
 - Catalog information on base tables
 - Selectivity (Reduction Factor) estimation
- **A search algorithm**
 - To sift through the plan space based on cost!



Summary

- **Query optimization is an important task in a relational DBMS.**
- **Must understand optimization in order to understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).**
- **Two parts to optimizing a query:**
 - Consider a set of alternative plans.
 - Must prune search space; typically, *left-deep plans* only.
 - Must estimate cost of each plan that is considered.
 - Must estimate size of result and cost for each plan node.
 - *Key issues:* Statistics, indexes, operator implementations.



Query Optimization

- **Query can be dramatically improved by changing access methods, order of operators.**
- **Iterator interface**
- **Cost estimation**
 - Size estimation and reduction factors
- **Statistics and Catalogs**
- **Relational Algebra Equivalences**
- **Choosing alternate plans**
- **Multiple relation queries**
- **Will focus on "System R"-style optimizers**



Highlights of System R Optimizer

- **Impact:**
 - Most widely used currently; works well for < 10 joins.
- **Cost estimation:**
 - Very inexact, but works ok in practice.
 - Statistics, maintained in system catalogs, used to estimate cost of operations and result sizes.
 - Considers combination of CPU and I/O costs.
 - More sophisticated techniques known now.
- **Plan Space: Too large, must be pruned.**
 - Many plans share common, "overpriced" subtrees
 - ignore them all!
 - In some implementations, only the space of *left-deep plans* is considered.
 - Cartesian products avoided in some implementations.



Query Blocks: Units of Optimization


- An SQL query is parsed into a collection of *query blocks*, and these are optimized one block at a time.
- Nested blocks are usually treated as calls to a subroutine, made once per outer tuple. (This is an over-simplification, wait til we learn more about nested queries.)

```
SELECT S.sname
FROM Sailors S
WHERE S.age IN
    (SELECT MAX (S2.age)
     FROM Sailors S2
     GROUP BY S2.rating)
```

Outer block Nested block

- ❖ For each block, the plans considered are:

- All available access methods, for each relation in FROM clause.
- All *left-deep join trees* (i.e., right branch always a base table, consider all join orders and join methods.)



Translating SQL to Relational Algebra

```
SELECT S.sid, MIN (R.day)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
GROUP BY S.sid
HAVING COUNT (*) >= 2
```

For each sailor with at least two reservations for red boats, find the sailor id and the earliest date on which the sailor has a reservation for a red boat.



Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)
Reserves (sid: integer, bid: integer, day: dates, rname: string)

- **Reserves:**

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages. 100 distinct bids.

- **Sailors:**

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
10 ratings, 40,000 sids.



Translating SQL to Relational Algebra

```
SELECT S.sid, MIN (R.day)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
GROUP BY S.sid
HAVING COUNT (*) >= 2
```

```

π S.sid, MIN(R.day)
(HAVING COUNT(*) > 2 (
  GROUP BY S.sid (
    σB.color = "red" (
      Sailors ⋈ Reserves ⋈ Boats))))

```



Relational Algebra Equivalences

- Allow us to choose different join orders and to 'push' selections and projections ahead of joins.

- **Selections:**

- $\sigma_{c_1 \wedge \dots \wedge c_n}(R) \equiv \sigma_{c_1}(\dots(\sigma_{c_n}(R))\dots)$ (*cascade*)
- $\sigma_{c_1}(\sigma_{c_2}(R)) \equiv \sigma_{c_1}(\sigma_{c_1}(R))$ (*commute*)

- **Projections:**

- $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1} \dots \pi_{an}(R))\dots)$ (*cascade*)

- **Cartesian Product**

$$- R \times (S \times T) \equiv (R \times S) \times T \quad (\text{associative})$$

$$-R \times S \equiv S \times R \quad (\text{commutative})$$

- This means we can do joins in any order.

- But...beware of cartesian product!



More Equivalences

- **Eager projection**

- Can cascade and "push" some projections thru selection
- Can cascade and "push" some projections below one side of a join
- Rule of thumb: can project anything not needed "downstream"

- **Selection between attributes of the two arguments of a cross-product converts cross-product to a join.**

- A selection on just attributes of R commutes with $R \bowtie S$. (i.e., $\sigma(R \bowtie S) \equiv \sigma(R) \bowtie S$)



Cost Estimation

- For each plan considered, must estimate total cost:
 - Must estimate *cost* of each operation in plan tree.
 - Depends on input cardinalities.
 - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
 - Must estimate *size of result* for each operation in tree!
 - Use information about the input relations.
 - For selections and joins, assume independence of predicates.
 - In System R, cost is boiled down to a single number consisting of $\#I/O + \text{factor} * \#CPU \text{ instructions}$
 - Q: Is "cost" the same as estimated "run time"?



Statistics and Catalogs

- Need information about the relations and indexes involved. *Catalogs* typically contain at least:
 - # tuples (**NTuples**) and # pages (**NPages**) per rel'n.
 - # distinct key values (**NKeys**) for each **index**.
 - low/high key values (**Low/High**) for each index.
 - Index height (**IHeight**) for each **tree** index.
 - # index pages (**INPages**) for each index.
- Catalogs updated periodically.
 - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- More detailed information (e.g., histograms of the values in some field) are sometimes stored.



Size Estimation and Reduction Factors

```
SELECT attribute list
FROM relation list
WHERE term1 AND ... AND termk
```

- Consider a query block:
- Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause.
- **Reduction factor (RF)** associated with each **term** reflects the impact of the **term** in reducing result size. **Result cardinality = Max # tuples * product of all RF's.**
- RF usually called "selectivity"
 - only R&G seem to call it Reduction Factor
 - beware of confusion between "high selectivity" as defined here and "highly selective" in common English!



Result Size Estimation

- Result cardinality = **Max # tuples * product of all RF's.**
- Term $col=value$ (given index I on col)

$$RF = 1/NKeys(I)$$
- Term $col1=col2$ (**This is handy for joins too...**)

$$RF = 1/MAX(NKeys(I1), NKeys(I2))$$
- Term $col > value$

$$RF = (High(I)-value)/(High(I)-Low(I))$$

(Implicit assumptions: values are uniformly distributed and terms are independent!)
- Note, if missing indexes, assume **1/10!!!**



Postgres 8: include/utills/selfuncs.h

```
/* default selectivity estimate
for equalities such as "A = b"
*/
#define DEFAULT_EQ_SEL 0.005

/* default selectivity estimate
for inequalities such as "A <
b" */
#define DEFAULT_INEQ_SEL
0.333333333333333333

/* default selectivity estimate
for range inequalities "A > b
AND A < c" */
#define DEFAULT_RANGE_INEQ_SEL
0.005

/* default selectivity estimate
for pattern-match operators
such as LIKE */
#define DEFAULT_MATCH_SEL 0.005

/* default number of distinct
values in a table */
#define DEFAULT_NUM_DISTINCT 200

/* default selectivity estimate
for boolean and null test
nodes */
#define DEFAULT_UNK_SEL
0.005

#define DEFAULT_NOT_UNK_SEL
(1.0 - DEFAULT_UNK_SEL)
```




Backend/optimizer/path/clausesel.c

```
/*
 *
 * THIS IS A HACK TO GET V4 OUT THE DOOR.
 * -- JMH 7/9/92
 */
s1 = (Selectivity) 0.3333333;
```



Think through estimation for joins

- Term $col1=col2$
 - $RF = 1/MAX(NKeys(I1), NKeys(I2))$
- Q: Given a join of R and S, what is the range of possible result sizes (in #of tuples)?**
 - If join is on a key for R (and a Foreign Key in S)?
 - A common case, can treat it specially
- General case: join on {A} ({A} is key for neither)**
 - estimate each tuple r of R generates $NTuples(S)/NKeys(A,S)$ result tuples, so...
 $NTuples(R) * NTuples(S)/NKeys(A,S)$
 - but can also consider it starting with S, yielding:
 $NTuples(S) * NTuples(R)/NKeys(A,R)$
 - If these two estimates differ, take the lower one!
 - Q: Why?



Enumeration of Alternative Plans

- There are two main cases:**
 - Single-relation plans
 - Multiple-relation plans
- For queries over a single relation, queries consist of a combination of selects, projects, and aggregate ops:**
 - Each available access path (file scan / index) is considered, and the one with the least estimated cost is chosen.
 - The different operations are essentially carried out together (e.g., if an index is used for a selection, projection is done for each retrieved tuple, and the resulting tuples are *pipelined* into the aggregate computation).



Cost Estimates for Single-Relation Plans

- Index I on primary key matches selection:**
 - Cost is $Height(I)+1$ for a B+ tree.
- Clustered index I matching one or more selects:**
 - $(NPages(I)+NPages(R)) * \text{product of RF's of matching selects.}$
- Non-clustered index I matching one or more selects:**
 - $(NPages(I)+NTuples(R)) * \text{product of RF's of matching selects.}$
- Sequential scan of file:**
 - $NPages(R)$.

☞ Recall: Must also charge for duplicate elimination if required



Example


```
SELECT S.sid
FROM Sailors S
WHERE S.rating=8
```

- If we have an index on rating:**
 - Cardinality = $(1/NKeys(I)) * NTuples(R) = (1/10) * 40000$ tuples
 - Clustered index: $(1/NKeys(I)) * (NPages(I)+NPages(R)) = (1/10) * (50+500) = 55$ pages are retrieved. (This is the *cost*.)
 - Unclustered index: $(1/NKeys(I)) * (NPages(I)+NTuples(R)) = (1/10) * (50+40000) = 401$ pages are retrieved.
- If we have an index on sid:**
 - Would have to retrieve all tuples/pages. With a clustered index, the cost is 50+500, with unclustered index, 50+40000.
- Doing a file scan:**
 - We retrieve all file pages (500).



Queries Over Multiple Relations

- A heuristic decision in System R:**
 - only left-deep join trees are considered.
 - As the number of joins increases, the number of alternative plans grows rapidly; we need to restrict the search space.
 - Left-deep trees allow us to generate all *fully pipelined* plans.
 - Intermediate results not written to temporary files.
 - Not all left-deep trees are fully pipelined (e.g., SM join).



Enumeration of Left-Deep Plans

- Left-deep plans differ only in the order of relations, the access method for each relation, and the join method for each join.**
- Enumerated using N passes (if N relations joined):**
 - Pass 1: Find best 1-relation plan for each relation.
 - Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (*All 2-relation plans.*)
 - Pass N: Find best way to join result of a (N-1)-relation plan (as outer) to the N'th relation. (*All N-relation plans.*)
- For each subset of relations, retain only:**
 - Cheapest plan overall, plus
 - Cheapest plan for each *interesting order* of the tuples.



The Dynamic Programming Table

Subset of tables in FROM clause	Interesting-order columns	Best plan	Cost
{R, S}	<none>	hashjoin(R, S)	1000
{R, S}	<R.a, S.b>	sortmerge(R, S)	1500



Enumeration of Plans (Contd.)

- An N-1 way plan is not combined with an additional relation unless there is a join condition between them, unless all predicates in WHERE have been used up.
 - i.e., avoid Cartesian products if possible.
- ORDER BY, GROUP BY, aggregates etc. handled as a final step, using either an 'interestingly ordered' plan or an additional sort/hash operator.
- In spite of pruning plan space, this approach is **still exponential** in the # of tables.
- Recall that in practice, COST considered is **#IOs + factor * CPU Inst**



Pass 1

- Best plan for accessing each relation regarded as the first relation in an execution plan
 - Reserves, Sailors: File Scan
 - Boats: B+ tree & Hash on color

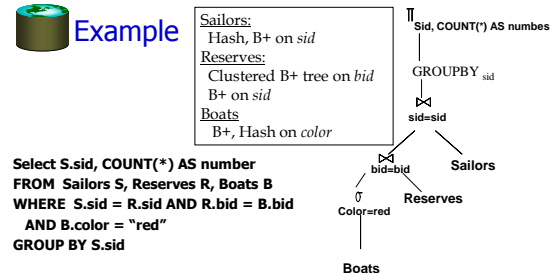


A Note on "Interesting Orders"

- An intermediate result has an "interesting order" if it is sorted by any of:
 - ORDER BY attributes
 - GROUP BY attributes
 - Join attributes of *yet-to-be-added* (downstream) joins



Example



- Pass1: Best plan(s) for accessing each relation
 - Reserves, Sailors: File Scan
 - Q: What about Clustered B+ on Reserves.bid???
 - Boats: B+ tree & Hash on color



Pass 2

- For each of the plans in pass 1, generate plans joining another relation as the inner, using all join methods (and matching inner access methods)
 - File Scan Reserves (outer) with Boats (inner)
 - File Scan Reserves (outer) with Sailors (inner)
 - File Scan Sailors (outer) with Boats (inner)
 - File Scan Sailors (outer) with Reserves (inner)
 - Boats hash on color with Sailors (inner)
 - Boats Btree on color with Sailors (inner)
 - Boats hash on color with Reserves (inner) (sort-merge)
 - Boats Btree on color with Reserves (inner) (BNL)
- Retain cheapest plan for each pair of relations



Pass 3 and beyond

- **For each of the plans retained from Pass 2, taken as the outer, generate plans for the next join**
 - eg Boats hash on color with Reserves (bid) (inner) (sortmerge))
 - inner Sailors (B-tree sid) sort-merge
- **Then, add the cost for doing the group by and aggregate:**
 - This is the cost to sort the result by sid, *unless it has already been sorted by a previous operator.*
- **Then, choose the cheapest plan**



Points to Remember

- **Must understand optimization in order to understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).**
- **Two parts to optimizing a query:**
 - Consider a set of alternative plans.
 - Good to prune search space; e.g., left-deep plans only, avoid Cartesian products.
 - Must estimate cost of each plan that is considered.
 - Output cardinality and cost for each plan node.
 - *Key issues:* Statistics, indexes, operator implementations.



Points to Remember

- **Single-relation queries:**
 - All access paths considered, cheapest is chosen.
 - *Issues:* Selections that *match* index, whether index key has all needed fields and/or provides tuples in a desired order.



More Points to Remember

- **Multiple-relation queries:**
 - All single-relation plans are first enumerated.
 - Selections/projections considered as early as possible.
 - Next, for each 1-relation plan, all ways of joining another relation (as inner) are considered.
 - Next, for each 2-relation plan that is `retained`, all ways of joining another relation (as inner) are considered, etc.
 - At each level, for each subset of relations, only best plan for each interesting order of tuples is `retained`.



Summary

- **Optimization is the reason for the lasting power of the relational system**
- **But it is primitive in some ways**
- **New areas: Smarter summary statistics (fancy histograms and "sketches"), auto-tuning statistics, adaptive runtime re-optimization (e.g. *eddies*)**