

SQL: The Query Language Part 2

R & G - Chapter 5

The important thing is not to stop questioning.

Albert Einstein



Example Database

Sailors

sid	sname	rating	age
22	Dustin	7	45
31	Lubber	8	55
95	Bob	3	63

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Reserves

sid	bid	day
22	101	10/10/06
95	103	11/12/06




Conceptual SQL Evaluation

```
SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualification
GROUP BY grouping-list
HAVING group-qualification
```

Project away columns
(just keep those used in
SELECT, GBY, HAVING)

Apply selections
(eliminate rows)

Relation
cross-product



Sorting the Results of a Query

- **ORDER BY** *column* [ASC | DESC] [, ...]

```
SELECT S.rating, S.sname, S.age
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid
      AND R.bid=B.bid AND B.color='red'
ORDER BY S.rating, S.sname;
```

- Can order by any column in SELECT list, including expressions or aggs:

```
SELECT S.sid, COUNT(*) AS redrescnt
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid
      AND R.bid=B.bid AND B.color='red'
GROUP BY S.sid
ORDER BY redrescnt DESC;
```



Null Values

- Field values in a tuple are sometimes **unknown** (e.g., a rating has not been assigned) or **inapplicable** (e.g., no spouse's name).
 - SQL provides a special value **null** for such situations.
- The presence of **null** complicates many issues. E.g.:
 - Special operators needed to check if value is/is not **null**.
 - Is $rating > 8$ true or false when $rating$ is equal to **null**? What about **AND**, **OR** and **NOT** connectives?
 - We need a **3-valued logic** (true, false and **unknown**).
 - Meaning of constructs must be defined carefully. (e.g., **WHERE** clause eliminates rows that don't evaluate to true.)
 - New operators (in particular, **outer joins**) possible/needed.



Joins

```
SELECT (column_list)
FROM table_name
[INNER | {LEFT | RIGHT | FULL} OUTER] JOIN table_name
ON qualification_list
WHERE ...
```

Explicit join semantics needed unless it is an INNER join (INNER is default)



Inner Join

Only the rows that match the search conditions are returned.

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

Returns only those sailors who have reserved boats
SQL-92 also allows:

```
SELECT s.sid, s.name, r.bid
FROM Sailors s NATURAL JOIN Reserves r
```

"NATURAL" means equi-join for each pair of attributes with the same name



```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103



Left Outer Join

Left Outer Join returns all matched rows, plus all unmatched rows from the table on the left of the join clause
(use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid
```

Returns all sailors & information on whether they have reserved boats



```
SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid
```

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

sid	bid	day
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	



Right Outer Join

Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT r.sid, b.bid, b.name
FROM Reserves r RIGHT OUTER JOIN Boats b
ON r.bid = b.bid
```

Returns all boats & information on which ones are reserved.



```
SELECT r.sid, b.bid, b.name
FROM Reserves r RIGHT OUTER JOIN Boats b
ON r.bid = b.bid
```

sid	bid	day
22	101	10/10/96
95	103	11/12/96

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine



Full Outer Join

Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.name
FROM Reserves r FULL OUTER JOIN Boats b
ON r.bid = b.bid
```

Returns all boats & all information on reservations



```
SELECT r.sid, b.bid, b.name
FROM Reserves r FULL OUTER JOIN Boats b
ON r.bid = b.bid
```

sid	bid	day	bid	bname	color
22	101	10/10/96	101	Interlake	blue
95	103	11/12/96	102	Interlake	red
			103	Clipper	green
			104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Note: in this case it is the same as the ROJ because bid is a foreign key in reserves, so all reservations must have a corresponding tuple in boats.



Views: Defining External DB Schemas

```
CREATE VIEW view_name
AS select_statement
```

Makes development simpler
Often used for security
Not "materialized"

```
CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scout
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```



Views Instead of Relations in Queries

```
CREATE VIEW Reds
AS SELECT B.bid, COUNT (*) AS scout
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

bid	scout	Reds
102	1	

```
SELECT bname, scout
FROM Reds R, Boats B
WHERE R.bid=B.bid
AND scout < 10
```



Discretionary Access Control

```
GRANT privileges ON object TO users
[WITH GRANT OPTION]
```

- Object can be a **Table** or a **View**
- Privileges can be:
 - Select
 - Insert
 - Delete
 - References (cols) – allow to create a foreign key that references the specified column(s)
 - All
- Can later be **REVOKED**
- Users can be single users or groups
- See Chapter 17 for more details.



Two more important topics

- **Constraints**
- **SQL embedded in other languages**



Integrity Constraints (Review)

- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- **Types of IC's:** Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints:* Field values must be of right type. Always enforced.
 - *Primary key and foreign key constraints:* you know them.



General Constraints

- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Checked on insert or update.
- Constraints can be named.

```
CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10 ))

CREATE TABLE Reserves
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK ('Interlake' <>
        (SELECT B.bname
         FROM Boats B
         WHERE B.bid=bid)))
```



Constraints Over Multiple Relations

- ```
CREATE TABLE Sailors
(sid INTEGER,
 sname CHAR(10),
 rating INTEGER,
 age REAL,
 PRIMARY KEY (sid),
 CHECK
 ((SELECT COUNT (S.sid) FROM Sailors S)
 + (SELECT COUNT (B.bid) FROM
 Boats B) < 100)

CREATE ASSERTION smallClub
CHECK
 ((SELECT COUNT (S.sid) FROM Sailors S)
 + (SELECT COUNT (B.bid)
 FROM Boats B) < 100)
```
- Awkward and wrong!
  - Only checks sailors!
  - Only required to hold if the associated table is non-empty.
  - ASSERTION is the right solution; not associated with either table.
  - Unfortunately, not supported in many DBMS.
  - Triggers are another solution.

*Number of boats  
plus number of  
sailors is < 100*



## Writing Applications with SQL

- SQL is not a general purpose programming language.

- + Tailored for data retrieval and manipulation
- + Relatively easy to optimize and parallelize
- Can't write entire apps in SQL alone

### Options:

Make the query language "Turing complete"

Avoids the "impedance mismatch"

but, loses advantages of relational language simplicity

Allow SQL to be embedded in regular programming languages.

Q: What needs to be solved to make the latter approach work?



## Embedded SQL

- DBMS vendors traditionally provided "host language bindings"
  - E.g. for C or COBOL
  - Allow SQL statements to be called from within a program
  - Typically you preprocess your programs
  - Preprocessor generates calls to a proprietary DB connectivity library
- General pattern
  - One call to *connect* to the right database (login, etc.)
  - SQL statements can refer to *host variables* from the language
- Typically vendor-specific
  - We won't look at any in detail, we'll look at standard stuff
- Problem
  - SQL relations are (multi-)sets, no *a priori* bound on the number of records. No such data structure in C.
  - SQL supports a mechanism called a *cursor* to handle this.



## Just to give you a flavor

```
EXEC SQL SELECT S.sname, S.age
INTO :c_sname, :c_age
FROM sailors S
WHERE S.sid = :c_sid
```



## Cursors

- Can declare a cursor on a relation or query
- Can *open* a cursor
- Can repeatedly *fetch* a tuple (moving the cursor)
- Special return value when all tuples have been retrieved.
- **ORDER BY** allows control over the order in which tuples are returned.
  - Fields in ORDER BY clause must also appear in SELECT clause.
- Can also modify/delete tuple pointed to by a cursor
  - A “non-relational” way to get a handle to a particular tuple
- There’s an Embedded SQL syntax for cursors
  - DECLARE <cursorname> CURSOR FOR <select stmt>
  - FETCH FROM <cursorname> INTO <variable names>
  - But we’ll peek at JDBC instead




## Database APIs: Alternative to embedding

- **Rather than modify compiler, add a library with database calls (API)**
  - special objects/methods
  - passes SQL strings from language, presents **result sets** in a language-friendly way
  - *ODBC* a C/C++ standard started on Windows
  - *JDBC* a Java equivalent
  - Most scripting languages have similar things
    - E.g. For Perl there is DBI, “oraPerl”, other packages
- **Mostly DBMS-neutral**
  - at least try to hide distinctions across different DBMSs



## Architecture



- A lookup service maps “data source names” (“DSNs”) to drivers
  - Typically handled by OS
- Based on the DSN used, a “driver” is linked into the app at runtime
- The driver traps calls, translates them into DBMS-specific code
- Database can be across a network
- ODBC is standard, so the same program can be used (in principle) to access multiple database systems
- Data source may not even be an SQL database!



## ODBC/JDBC

- **Various vendors provide drivers**
  - MS bundles a bunch into Windows
  - Vendors like DataDirect and OpenLink sell drivers for multiple OSes
- **Drivers for various data sources**
  - Relational DBMSs (Oracle, DB2, SQL Server, etc.)
  - “Desktop” DBMSs (Access, dBase, Paradox, FoxPro, etc.)
  - Spreadsheets (MS Excel, Lotus 1-2-3, etc.)
  - Delimited text files (.CSV, .TXT, etc.)
- **You can use JDBC/ODBC clients over many data sources**
  - E.g. MS Query comes with many versions of MS Office (msqry32.exe)
- **Can write your own Java or C++ programs against xDBC**



## JDBC

- **Part of Java, very easy to use**
- **Java comes with a JDBC-to-ODBC bridge**
  - So JDBC code can talk to any ODBC data source
  - E.g. look in your Windows Control Panel or MacOS Utilities folder for JDBC/ODBC drivers!
- **JDBC tutorial online**
  - <http://developer.java.sun.com/developer/Books/JDBC/Tutorial/>



## Ruby on Rails

- **Rails’ find method gives a simple rowset interface**
  - Just an array of records
  - Unfortunately slurps entire result set into memory.
- **Rails’ ORM (Object Relational Mapping) goes beyond queries and cursors**
  - Data modeling and implicit query construction
  - The ActiveRecord.find method sometimes generates key/foreign-key joins, for example
- **Can also do:**
  - find\_by\_sql
  - ActiveRecordBase.connection.execute/



## API Summary

**APIs are needed to interface DBMSs to programming languages**

- Embedded SQL uses "native drivers" and is usually faster but less standard
- ODBC (used to be Microsoft-specific) for C/C++
- JDBC the standard for Java
- Scripting languages (PHP, Perl, JSP) are becoming the preferred technique for web-based systems