


Physical Database Design and Tuning

R&G - Chapter 20

Although the whole of this life were said to be nothing but a dream and the physical world nothing but a phantasm, I should call this dream or phantasm real enough, if, using reason well, we were never deceived by it.

Baron Gottfried Wilhelm von Leibniz



Understanding the Workload

- **For each query in the workload:**
 - Which relations does it access?
 - Which attributes are retrieved?
 - Which attributes are involved in selection/join conditions?
How selective are these conditions likely to be?
- **For each update in the workload:**
 - Which attributes are involved in selection/join conditions?
How selective are these conditions likely to be?
 - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected.



Decisions to Make

- **What indexes should we create?**
 - Which relations should have indexes? What field(s) should be the search key? Should we build several indexes?
- **For each index, what kind of an index should it be?**
 - Clustered? Dynamic/static?
- **Should we make changes to the conceptual schema?**
 - For example, denormalize
- **Horizontal partitioning, replication, views ...**



Introduction

- **We have talked at length about "database design"**
 - Conceptual Schema: info to capture, tables, columns, views, etc.
 - Physical Schema: indexes, clustering, etc.
- **Physical design linked tightly to query optimization**
- **We must begin by understanding the *workload*:**
 - The most important queries and how often they arise.
 - The most important updates and how often they arise.
 - The desired performance for these queries and updates.



Creating an ISUD Chart

Insert, Select, Update, Delete Frequencies

Transaction	Frequency	% table	Employee Table		
			Name	Salary	Address
Payroll Run	monthly	100	S	S	S
Add Emps	daily	0.1	I	I	I
Delete Emps	daily	0.1	D	D	D
Give Raises	monthly	10	S	U	



Index Selection

- **One approach:**
 - Consider most important queries in turn.
 - Consider best plan using the current indexes, and see if better plan is possible with an additional index.
 - If so, create it.
- **Before creating an index, must also consider the impact on updates in the workload!**
 - Trade-off: indexes can make queries go faster, updates slower. Require disk space, too.



Example 1

```
SELECT E.ename, D.mgr
FROM Emp E, Dept D
WHERE E.dno=D.dno AND D.dname='Toy'
```

- **B+ tree index on *D.dname* supports 'Toy' selection.**
 - Given this, index on *D.dno* is not needed.
- **B+ tree index on *E.dno* allows us to get matching (inner) Emp tuples for each selected (outer) Dept tuple.**
- **What if WHERE included: `` ... AND *E.age*=25'' ?**
 - Could retrieve Emp tuples using index on *E.age*, then join with Dept tuples satisfying *dname* selection. Comparable to strategy that used *E.dno* index.
 - So, if *E.age* index is already created, this query provides much less motivation for adding an *E.dno* index.



Example 2

```
SELECT E.ename, D.mgr
FROM Emp E, Dept D
WHERE E.sal BETWEEN 10000 AND 20000
AND E.hobby='Stamps' AND E.dno=D.dno
```

- **All selections are on Emp so it should be the outer relation in any Index NL join.**
 - Suggests that we build a B+ tree index on *D.dno*.
- **What index should we build on Emp?**
 - B+ tree on *E.sal* could be used, OR an index on *E.hobby* could be used. Only one of these is needed, and which is better depends upon the selectivity of the conditions.
 - As a rule of thumb, equality selections more selective than range selections.
- **As both examples indicate, our choice of indexes is guided by the plan(s) that we expect an optimizer to consider for a query. *Have to understand optimizers!***



Examples of Clustering

- **B+ tree index on *E.age* can be used to get qualifying tuples.**
 - How selective is the condition?
 - Is the index clustered?
- **Consider the GROUP BY query.**
 - If many tuples have *E.age* > 10, using *E.age* index and sorting the retrieved tuples may be costly.
 - Clustered *E.dno* index may be better!
- **Equality queries and duplicates:**
 - Clustering on *E.hobby* helps!

```
SELECT E.dno
FROM Emp E
WHERE E.age>40
```

```
SELECT E.dno, COUNT(*)
FROM Emp E
WHERE E.age>10
GROUP BY E.dno
```

```
SELECT E.dno
FROM Emp E
WHERE E.hobby=Stamps
```



Index-Only Plans

- **A number of queries can be answered without retrieving any tuples from one or more of the relations involved if a suitable index is available.**

<*E.dno*>
<*E.dno, E.eid*>
<*E.dno*>
<*E.dno, E.sal*>
B-tree trick!

```
SELECT D.mgr
FROM Dept D, Emp E
WHERE D.dno=E.dno
```

```
SELECT E.dno, COUNT(*)
FROM Emp E
GROUP BY E.dno
```

```
SELECT E.dno, MIN(E.sal)
FROM Emp E
GROUP BY E.dno
```

```
SELECT AVG(E.sal)
FROM Emp E
WHERE E.age=25 AND
E.sal BETWEEN 3000 AND 5000
```

<*E.age, E.sal*>
or
<*E.sal, E.age*>



Horizontal Decompositions

- **Usual Def. of decomposition: Relation is replaced by collection of relations that are *projections*. Most important case.**
 - We will talk about this at length as part of Conceptual DB Design
- **Sometimes, might want to replace relation by a collection of relations that are *selections*.**
 - Each new relation has same schema as original, but subset of rows.
 - Collectively, new relations contain all rows of the original.
 - Typically, the new relations are disjoint.



Horizontal Decompositions (Contd.)

- **Contracts (*Cid*, *Sid*, *Jid*, *Did*, *Pid*, *Qty*, *Val*)**
- **Suppose that contracts with value > 10000 are subject to different rules.**
 - So queries on Contracts will often say *WHERE val>10000*.
- **One approach: clustered B+ tree index on the *val* field.**
- **Second approach: replace contracts by two new relations, *LargeContracts* and *SmallContracts*, with the same attributes (CSJDPQV).**
 - Performs like index on such queries, but no index overhead.
 - Can build clustered indexes on other attributes, in addition!



Masking Conceptual Schema Changes

```
CREATE VIEW Contracts(cid, sid, jid, did, pid, qty, val)
AS SELECT *
FROM LargeContracts
UNION
SELECT *
FROM SmallContracts
```

- **Horizontal Decomposition** from above
- **Masked by a view.**
 - NOTE: queries with condition *val* > 10000 must be asked wrt LargeContracts for efficiency: so some users may have to be aware of change.
 - I.e. the users who were having performance problems
 - Arguably that's OK -- they wanted a solution!



Tuning Queries and Views

- If a query runs slower than expected, check if an index needs to be re-clustered, or if statistics are too old.
- Sometimes, the DBMS may not be executing the plan you had in mind. Common areas of weakness:
 - Selections involving null values (bad selectivity estimates)
 - Selections involving arithmetic or string expressions (ditto)
 - Selections involving OR conditions (ditto)
 - Complex subqueries (more on this later)
 - Lack of evaluation features like index-only strategies or certain join methods or poor size estimation.
- Check the plan that is being used! Then adjust the choice of indexes or rewrite the query/view.
 - E.g. check via POSTGRES "Explain" command
 - Some systems rewrite for you under the covers (e.g. DB2)
 - Can be confusing and/or helpful!



Index Tuning "Wizards"

- Both IBM's DB2 and MS SQL Server have automated index advisors
 - Some info in Section 20.6 of the book
- **Basic idea:**
 - They take a workload of queries
 - Possibly based on logging what's been going on
 - They use the optimizer cost metrics to estimate the cost of the workload over different choices of sets of indexes
 - Enormous # of different choices of sets of indexes:
 - Heuristics to help this go faster



More Guidelines for Query Tuning

- Minimize the use of DISTINCT: don't need it if duplicates are acceptable, or if answer contains a key.
- Minimize the use of GROUP BY and HAVING:

```
SELECT MIN(E.age)
FROM Employee E
GROUP BY E.dno
HAVING E.dno=102
```

```
SELECT MIN(E.age)
FROM Employee E
WHERE E.dno=102
```

- Consider DBMS use of index when writing arithmetic expressions: $E.age = 2 * D.age$ will benefit from index on $E.age$, but might not benefit from index on $D.age$!



Guidelines for Query Tuning (Contd.)

- Avoid using intermediate relations:

vs.

```
SELECT E.dno, AVG(E.sal)
FROM Emp E, Dept D
WHERE E.dno=D.dno
AND D.mgrname='Joe'
GROUP BY E.dno
```

```
SELECT * INTO Temp
FROM Emp E, Dept D
WHERE E.dno=D.dno
AND D.mgrname='Joe'
```

and

```
SELECT T.dno, AVG(T.sal)
FROM Temp T
GROUP BY T.dno
```

- Does not materialize the intermediate reln Temp.
- If there is a dense B+ tree index on $\langle dno, sal \rangle$, an index-only plan can be used to avoid retrieving Emp tuples in the second query!



Points to Remember

- Indexes must be chosen to speed up important queries (and perhaps some updates!).
 - Index maintenance overhead on updates to key fields.
 - Choose indexes that can help many queries, if possible.
 - Build indexes to support index-only strategies.
 - Clustering is an important decision; only one index on a given relation can be clustered!
 - Order of fields in composite index key can be important.
- Static indexes may have to be periodically re-built.
- Statistics have to be periodically updated.



Points to remember (Contd.)

- **Over time, indexes have to be fine-tuned (dropped, created, re-clustered, ...) for performance.**
 - Should determine the plan used by the system, and adjust the choice of indexes appropriately.
- **System may still not find a good plan:**
 - Only left-deep plans?
 - Null values, arithmetic conditions, string expressions, the use of ORs, nested queries, etc. can confuse an optimizer.
- **So, may have to rewrite the query/view:**
 - Avoid nested queries, temporary relations, complex conditions, and operations like DISTINCT and GROUP BY.