


Concurrency Control

R&G - Chapter 17



Smile, it is the key that fits the lock of everybody's heart.

Anthony J. D'Angelo,
The College Blue Book



Review

- ACID transaction semantics.
- Today: focus on Isolation property
 - Serial schedules safe but slow
 - Try to find schedules **equivalent** to serial ...



Conflicting Operations

- Need a tool to decide if 2 schedules are equivalent
- Use notion of "conflicting" operations
- Definition: Two operations **conflict** if:
 - They are by different transactions,
 - they are on the same object,
 - and at least one of them is a write.



Conflict Serializable Schedules

- Definition: Two schedules are **conflict equivalent** iff:
 - They involve the same actions of the same transactions, and
 - every pair of conflicting actions is ordered the same way
- Definition: Schedule S is **conflict serializable** if:
 - S is conflict equivalent to some serial schedule.
- Note, some "serializable" schedules are NOT conflict serializable
 - A price we pay to achieve efficient enforcement.



Conflict Serializability – Intuition

- A schedule S is conflict serializable if:
 - You are able to transform S into a serial schedule by swapping **consecutive non-conflicting** operations of different transactions.
- *Example*:

$$\begin{array}{ccccc}
 R(A) & W(A) & & R(B) & W(B) \\
 & & R(A) & W(A) & \\
 & & & & R(B) & W(B) \\
 & & & & & \\
 & & & & & \equiv \\
 R(A) & W(A) & R(B) & W(B) & & \\
 & & & & R(A) & W(A) & R(B) & W(B)
 \end{array}$$


Conflict Serializability (Continued)

- Here's another example:

$$\begin{array}{ccc}
 R(A) & & W(A) \\
 & R(A) & W(A)
 \end{array}$$

- Serializable or not????

NOT!



Dependency Graph




- **Dependency graph:**
 - One node per Xact
 - Edge from Ti to Tj if:
 - An operation Oi of Ti conflicts with an operation Oj of Tj and
 - Oi appears earlier in the schedule than Oj.
- **Theorem:** Schedule is conflict serializable *if and only if* its dependency graph is acyclic.



Example

- A schedule that is not conflict serializable:

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	



- The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.



An Aside: View Serializability

- Alternative (weaker) notion of serializability.
- Schedules S1 and S2 are **view equivalent** if:
 1. If Ti reads initial value of A in S1, then Ti also reads initial value of A in S2
 2. If Ti reads value of A written by Tj in S1, then Ti also reads value of A written by Tj in S2
 3. If Ti writes final value of A in S1, then Ti also writes final value of A in S2
- Basically, allows all conflict serializable schedules + "blind writes"

T1: R(A)	W(A)	
T2: W(A)		
T3: W(A)		

view

T1: R(A), W(A)		
T2: W(A)		
T3: W(A)		



Notes on Serializability Definitions

- View Serializability allows (slightly) more schedules than Conflict Serializability does.
 - Problem is that it is difficult to enforce efficiently.
- Neither definition allows all schedules that you would consider "serializable".
 - This is because they don't understand the meanings of the operations or the data.
- In practice, Conflict Serializability is what gets used, because it can be enforced efficiently.
 - To allow more concurrency, some special cases do get handled separately, such as for travel reservations, etc.



Two-Phase Locking (2PL)

Lock
Compatibility
Matrix


	S	X
S	✓	–
X	–	–

rules:

- Xact must obtain a **S** (*shared*) lock before reading, and an **X** (*exclusive*) lock before writing.
- Xact cannot get new locks after releasing any locks.



Two-Phase Locking (2PL), cont.



2PL guarantees conflict serializability



But, does not prevent **Cascading Aborts**.





Strict 2PL


- *Problem:* Cascading Aborts
- *Example:* rollback of T1 requires rollback of T2!

T1:	R(A), W(A),	R(B), W(B), Abort
T2:	R(A), W(A)	

- Strict Two-phase Locking (Strict 2PL) protocol:
Same as 2PL, except:
Locks released only when transaction completes
i.e., either:
(a) transaction has committed (commit record on disk),
or
(b) transaction has aborted and rollback is complete.



Strict 2PL (continued)



Next ...

- A few examples



Non-2PL, A= 1000, B=2000, Output =?

Lock_X(A)	
Read(A)	Lock_S(A)
A := A-50	
Write(A)	
Unlock(A)	
	Read(A)
	Unlock(A)
	Lock_S(B)
Lock_X(B)	
	Read(B)
	Unlock(B)
	PRINT(A+B)
Read(B)	
B := B +50	
Write(B)	
Unlock(B)	



2PL, A= 1000, B=2000, Output =?

Lock_X(A)	
Read(A)	Lock_S(A)
A := A-50	
Write(A)	
Lock_X(B)	
Unlock(A)	
	Read(A)
	Lock_S(B)
Read(B)	
B := B +50	
Write(B)	
Unlock(B)	Unlock(A)
	Read(B)
	Unlock(B)
	PRINT(A+B)

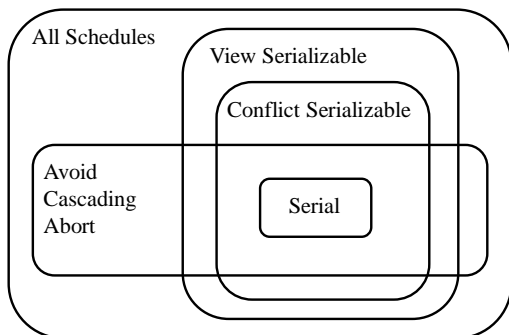


Strict 2PL, A= 1000, B=2000, Output =?


Lock_X(A)	
Read(A)	Lock_S(A)
A := A-50	
Write(A)	
Lock_X(B)	
Read(B)	
B := B +50	
Write(B)	
Unlock(A)	
Unlock(B)	
	Read(A)
	Lock_S(B)
	Read(B)
	PRINT(A+B)
	Unlock(A)
	Unlock(B)



Venn Diagram for Schedules



Which schedules does Strict 2PL allow?



Lock Management

- Lock and unlock requests handled by **Lock Manager**
- LM keeps an entry for each currently held lock.
- Entry contains:
 - List of xacts currently holding lock
 - Type of lock held (shared or exclusive)
 - Queue of lock requests



Lock Management, cont.

- When lock request arrives:
 - Does any other xact hold a conflicting lock?
 - If no, grant the lock.
 - If yes, put requestor into wait queue.
- Lock upgrade:
 - xact with shared lock can request to upgrade to exclusive



Example

Lock_X(A)	
	Lock_S(B)
	Read(B)
	Lock_S(A)
Read(A)	
A: = A-50	
Write(A)	
Lock_X(B)	



Deadlocks

- **Deadlock:** Cycle of transactions waiting for locks to be released by each other.
- Two ways of dealing with deadlocks:
 - prevention
 - detection
- Many systems just punt and use Timeouts
 - What are the dangers with this approach?



Deadlock Detection


- Create and maintain a “waits-for” graph
- Periodically check for cycles in graph



Deadlock Detection (Continued)

Example:

T1: S(A), S(D),
 T2: X(B)
 T3: S(D), S(C),
 T4: X(B)



Deadlock Prevention

- Assign priorities based on timestamps.
- Say T_i wants a lock that T_j holds
 Two policies are possible:
Wait-Die: If T_i has higher priority, T_i waits for T_j ; otherwise T_i aborts
Wound-wait: If T_i has higher priority, T_j aborts; otherwise T_i waits
- Why do these schemes guarantee no deadlocks?
- Important detail: If a transaction re-starts, make sure it gets its original timestamp. -- Why?



Summary

- Correctness criterion for isolation is “serializability”.
 – In practice, we use “conflict serializability,” which is somewhat more restrictive but easy to enforce.
- Two Phase Locking and Strict 2PL: Locks implement the notions of conflict directly.
 – The lock manager keeps track of the locks issued.
 – **Deadlocks** may arise; can either be prevented or detected.